

Lecture 1

State Space Models and the Kalman Filter

Lars A. Lochstoer

UCLA Anderson School of Management

Outline

① Latent (or Hidden) processes

- ▶ Learning from a noisy signal

② Sequential Learning

③ The Local Level Model

④ State space models and the Kalman Filter

- ▶ Filtering Probabilities
- ▶ Smoothing Probabilities
- ▶ Maximum Likelihood Estimation

Latent Processes

Economy driven by forces we do not directly observe:

- ① Risk aversion / risk capacity
- ② The state of the business cycle
- ③ Beliefs about asset payoffs, macroeconomic conditions, etc.
- ④ Monetary and fiscal stance (e.g., true policy parameters)
- ⑤ +++

These are generally time-varying, dynamic concepts – not constant parameters

We generally do observe outcomes that are related to these latent (hidden) drivers

- We can learn from such (noisy) signals!

Noisy signals: a simple example

Assume inflation π_t follows the process:

$$\pi_t = \mu + \phi(\pi_{t-1} - \mu) + x_t + \varepsilon_t,$$

where $\varepsilon_t \sim N(0, \sigma^2)$ and x_t is a latent variable.

Do we observe a (noisy) signal about x_t ? Yes!

- Define

$$z_t \equiv \pi_t - \mu - \phi(\pi_{t-1} - \mu).$$

- Then

$$z_t = x_t + \varepsilon_t, \quad \text{and} \quad z_t | x_t \sim N(x_t, \sigma^2).$$

Given what we have so far, what is $E_t(x_t)$?

$$E_t(x_t) = z_t.$$

Aside: The expectation operator and latent variables

The expectation operator

$$E_t(\cdot) = E(\cdot | I_t)$$

refers to the expectation conditional on the *information set* at time t , I_t .

The latent variable x_t is not in your information set at time t , despite it having a t subscript.

What is in our information set then?

- The history of realized inflation π , and thus signals z , up until time t
- This history is typically written π^t and z^t

Noisy signals and prior beliefs

The above expectation

$$E_t(x_t) = z_t$$

assumed we had no prior knowledge about what x_t might be.

- Equivalently, we have a “flat” prior about x_t

Priors and posteriors? Let's revisit Bayes rule!

Bayes Rule and Updates in Beliefs

Recall Bayes rule:

$$P(A|B)P(B) = P(A, B) \Leftrightarrow P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

So, in our example:

$$P(x_t|z_t) = \frac{P(z_t|x_t)P(x_t)}{P(z_t)}$$

The lingo:

- $P(x_t)$: the **prior distribution** of x_t
- $P(x_t|z_t)$: the **posterior distribution** after observing the signal
- $P(z_t|x_t)$: the pdf of the signal conditional on x_t : the **likelihood function**
- $P(z_t)$: the marginal dist. of the observation z_t (thus, $P(z_t)$ is a constant)

Bayes Rule and Updates in Beliefs (cont'd)

Note that in our case:

$$P(z_t | x_t) = N(x_t, \sigma^2).$$

Let the prior be Normally distributed:

$$p(x_t) = N(x_{t|t-1}, \sigma_{t|t-1}^2)$$

The appendix shows that in this case

$$\frac{P(z_t | x_t) P(x_t)}{P(z_t)} = N(x_{t|t}, \sigma_{t|t}^2),$$

where

$$x_{t|t} = \frac{\sigma^{-2}}{\sigma^{-2} + \sigma_x^{-2}} z_t + \frac{\sigma_x^{-2}}{\sigma^{-2} + \sigma_x^{-2}} x_{t|t-1},$$

and

$$\sigma_{t|t}^2 = \left(\sigma_{t|t-1}^{-2} + \sigma^{-2} \right)^{-1}.$$

Flat vs informative priors

A flat prior means $\sigma_x^2 = \infty$

- We then have:

$$x_{t|t} = \frac{\sigma^{-2}}{\sigma^{-2} + \sigma_{t|t-1}^{-2}} z_t + \frac{\sigma_x^{-2}}{\sigma^{-2} + \sigma_{t|t-1}^{-2}} x_{t|t-1} = z_t.$$

But, if we have an informative prior, $\sigma_x^2 < \infty$

- Such an informative prior could come from prior observations if x_t is not independent across time
- For instance, x_t might follow an AR(1) process

Sequential Learning

We update beliefs every time new information arrives

- This is called **sequential learning**

Let y^t denote available data up until and including time t

- Assume we are trying to infer the value of a latent variable, x_t . If the data only provides noisy signals, we are looking for the distribution of $x_t|y^t$.
- Note that $x_t|y^t = x_t|y_t, y^{t-1}$

Sequential learning and Bayes Rule:

$$p(x_t|y^t) = p(x_t|y_t, y^{t-1}) \propto p(y_t|x_t, y^{t-1}) p(x_t|y^{t-1})$$

- $p(x_t|y^{t-1})$ is the *prior distribution* of the state x_t , conditional on data up until time $t - 1$
- $p(y_t|x_t, y^{t-1})$ is the *likelihood* of the empirical observation y_t , given your prior
- $p(x_t|y^t)$ is the *posterior distribution* about the latent state, x_t

The Local Level Model

Local Level Model

Consider the **local level model**, for $t = 1, \dots, T$:

$$\begin{aligned}y_t &= \mu_t + \eta_t, & \eta_t &\sim N(0, \sigma_\eta^2) \\ \mu_{t+1} &= \mu_t + \varepsilon_{t+1}, & \varepsilon_{t+1} &\sim N(0, \sigma_\varepsilon^2)\end{aligned}$$

- We observe the data y_t for $t = 1, \dots, T$
- But! The trend μ_t is not observable. It is latent.
- The Kalman filter provides a way to estimate the trend μ_t .
- The trend μ_t is also known as a **state variable**.
- In this lecture, I will use the notation: $y_{1:t} = (y_1, \dots, y_t)$

Local Level Model

$$y_t = \mu_t + \eta_t, \quad \eta_t \sim N(0, \sigma_\eta^2)$$
$$\mu_{t+1} = \mu_t + \varepsilon_{t+1}, \quad \varepsilon_{t+1} \sim N(0, \sigma_\varepsilon^2)$$

- The Kalman filter, Kalman (1960), is a recursive algorithm that performs:
 - ① Filtering
 - ② Prediction
 - ③ Smoothing
 - ④ Evaluation of the log-likelihood
- Let's introduce each of these ideas one-by-one.
- For now, assume we know the parameters: $\theta = (\sigma_\eta^2, \sigma_\varepsilon^2)$

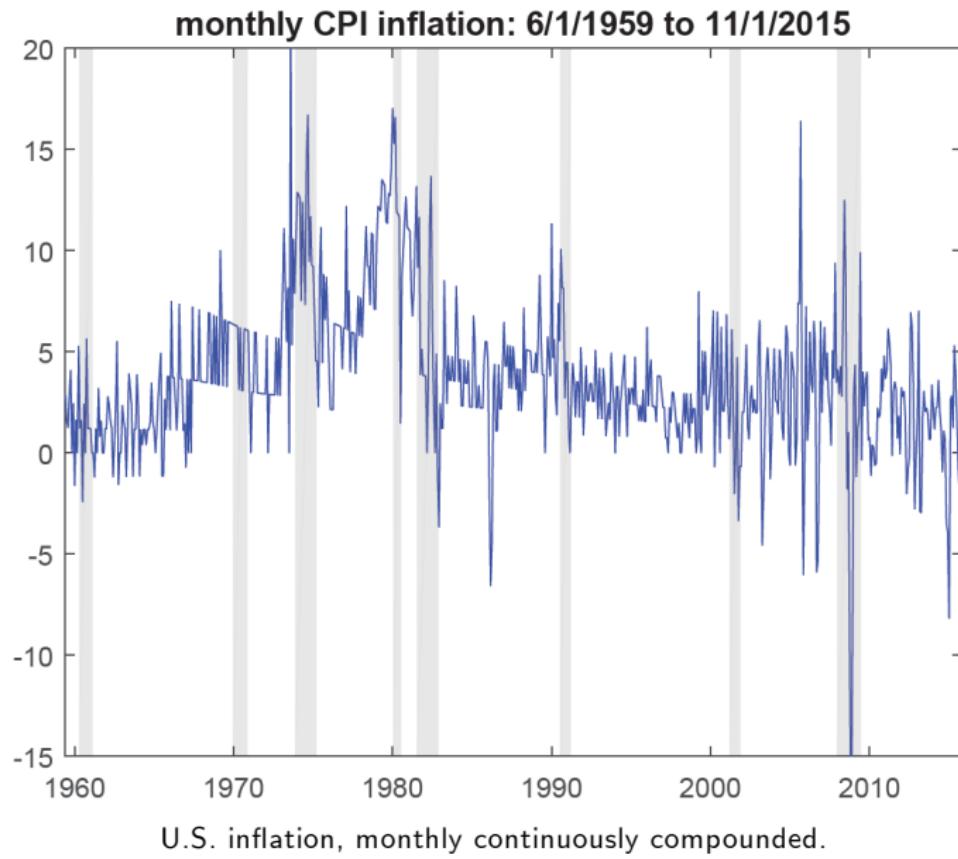
U.S. inflation

- Let $y_t = \pi_t$ denote monthly U.S. CPI inflation
- Stock and Watson (2007) use the local level model to forecast inflation.

$$\begin{aligned}\pi_t &= \bar{\pi}_t + \eta_t, & \eta_t &\sim N(0, \sigma_\eta^2) \\ \bar{\pi}_{t+1} &= \bar{\pi}_t + \varepsilon_{t+1}, & \varepsilon_{t+1} &\sim N(0, \sigma_\varepsilon^2)\end{aligned}$$

- The state variable $\mu_t = \bar{\pi}_t$ represents the trend in inflation and is what many economists would call '*expected inflation*'.

U.S. inflation



Bayes Theorem

- Let's start at time $t = 0$ before we have seen the first observation y_1 .
- Let $p(\mu_1; \theta)$ denote a prior distribution that describes our beliefs about μ_1 before we have observed y_1 .
- Let's assume that our prior $p(\mu_1; \theta)$ is a normal distribution

$$p(\mu_1; \theta) = N(\mu_{1|0}, \Sigma_{1|0})$$

- The notation $\mu_{1|0}$ means that it is our guess (prediction) of the value of μ_1 at time $t = 1$ but with time $t = 0$ information.
- The variance $\Sigma_{1|0}$ measures our uncertainty about μ_1 with time $t = 0$ information.
- Recall that, for the normal distribution, the mean and covariance matrix are **sufficient statistics**.

Bayes Theorem

- We have the prior: $p(\mu_1; \theta) = N(\mu_{1|0}, \Sigma_{1|0})$.
- We observe y_1 which is a noisy measure of μ_1 .
- Our model says y_1 and μ_1 are related:

$$y_1 = \mu_1 + \eta_1, \quad \eta_1 \sim N(0, \sigma_\eta^2)$$

- The (conditional) likelihood is:

$$p(y_1 | \mu_1; \theta) = N(\mu_1, \sigma_\eta^2)$$

Bayes Theorem

- We have the prior: $p(\mu_1; \theta) = N(\mu_{1|0}, \Sigma_{1|0})$
- The (conditional) likelihood: $p(y_1 | \mu_1; \theta) = N(\mu_1, \sigma_\eta^2)$
- After we observe the data y_1 , how do we revise our beliefs about μ_1 ?
- We apply Bayes rule:

$$p(\mu_1 | y_1; \theta) = \frac{p(y_1 | \mu_1; \theta)p(\mu_1; \theta)}{p(y_1; \theta)}$$

- The posterior $p(\mu_1 | y_1; \theta)$ describes our beliefs about μ_1 after we observe y_1 .

Filtering: local level model

- Applying Bayes Rule, we find

$$p(\mu_1 | y_1; \theta) = \frac{p(y_1 | \mu_1; \theta)p(\mu_1; \theta)}{p(y_1; \theta)} = N(\mu_{1|1}, \Sigma_{1|1})$$

where the mean and variance are

$$\mu_{1|1} = \mu_{1|0} + \Sigma_{1|0} F_1^{-1} (y_1 - \mu_{1|0}) \quad \Sigma_{1|1} = \Sigma_{1|0} - \Sigma_{1|0} F_1^{-1} \Sigma_{1|0}$$

and $F_1 = \Sigma_{1|0} + \sigma_\eta^2$.

- The notation $\mu_{1|1}$ indicates our estimate for μ_1 given $t = 1$ information.
- All distributions are normally distributed. Only update the mean and covariance matrix (the sufficient statistics).

Aside: Updating equations versus earlier case

Updating equations look superficially different from those earlier in this note. But:

$$\begin{aligned}\mu_{1|1} &= \mu_{1|0} + \Sigma_{1|0} F_1^{-1} (y_1 - \mu_{1|0}) \\&= \mu_{1|0} \left(1 - \Sigma_{1|0} F_1^{-1} \right) + \Sigma_{1|0} F_1^{-1} y_1 \\&= \mu_{1|0} \left(1 - \frac{\Sigma_{1|0}}{\Sigma_{1|0} + \sigma_\eta^2} \right) + \frac{\Sigma_{1|0}}{\Sigma_{1|0} + \sigma_\eta^2} y_1 \\&= \frac{\sigma_\eta^2}{\Sigma_{1|0} + \sigma_\eta^2} \mu_{1|0} + \frac{\Sigma_{1|0}}{\Sigma_{1|0} + \sigma_\eta^2} y_1 \\&= \frac{\sigma_\eta^2}{\Sigma_{1|0} + \sigma_\eta^2} \frac{\Sigma_{1|0}^{-1} \sigma_\eta^{-2}}{\Sigma_{1|0}^{-1} \sigma_\eta^{-2}} \mu_{1|0} + \frac{\Sigma_{1|0}}{\Sigma_{1|0} + \sigma_\eta^2} \frac{\Sigma_{1|0}^{-1} \sigma_\eta^{-2}}{\Sigma_{1|0}^{-1} \sigma_\eta^{-2}} y_1 \\&= \frac{\Sigma_{1|0}^{-1}}{\Sigma_{1|0}^{-1} + \sigma_\eta^{-2}} \mu_{1|0} + \frac{\sigma_\eta^{-2}}{\Sigma_{1|0}^{-1} + \sigma_\eta^{-2}} y_1,\end{aligned}$$

which is of the same form as the initial updating example.

Aside: Updating equations versus earlier case

Next, the variance:

$$\begin{aligned}\Sigma_{1|1} &= \Sigma_{1|0} - \Sigma_{1|0} F_1^{-1} \Sigma_{1|0} \\&= \Sigma_{1|0} - \frac{\Sigma_{1|0}^2}{\Sigma_{1|0} + \sigma_\eta^2} \\&= \frac{\Sigma_{1|0}^2 + \Sigma_{1|0} \sigma_\eta^2 - \Sigma_{1|0}^2}{\Sigma_{1|0} + \sigma_\eta^2} \\&= \frac{\Sigma_{1|0} \sigma_\eta^2}{\Sigma_{1|0} + \sigma_\eta^2} \frac{\Sigma_{1|0}^{-1} \sigma_\eta^{-2}}{\Sigma_{1|0}^{-1} \sigma_\eta^{-2}} \\&= \frac{1}{\Sigma_{1|0}^{-1} + \sigma_\eta^{-2}},\end{aligned}$$

which is of the same form as the initial updating example.

One-step Ahead Prediction

- Suppose we want to forecast the value of μ_2 given data up to y_1 .
- We want the one-step ahead predictive distribution: $p(\mu_2|y_1; \theta)$
- The predictive distribution $p(\mu_2|y_1; \theta)$ describes our uncertainty about μ_2 given we observe y_1 .
- Our model says μ_2 and μ_1 are related:

$$\mu_2 = \mu_1 + \varepsilon_2, \quad \varepsilon_2 \sim N(0, \sigma_\varepsilon^2)$$

- This defines the (Markov) transition density

$$p(\mu_2|\mu_1; \theta) = N(\mu_1, \sigma_\varepsilon^2)$$

One-step Ahead Prediction: local level model

- To calculate $p(\mu_2|y_1; \theta)$, we integrate out μ_1 by

$$p(\mu_2|y_1; \theta) = \int p(\mu_2|\mu_1; \theta)p(\mu_1|y_1; \theta)d\mu_1$$

- Since all distributions are Gaussian, integral can be solved analytically

$$p(\mu_2|y_1; \theta) = N(\mu_{2|1}, \Sigma_{2|1})$$

where the mean and variance are

$$\mu_{2|1} = \mu_{1|1} \quad \Sigma_{2|1} = \Sigma_{1|1} + \sigma_\varepsilon^2$$

- We just update the sufficient statistics.
- At time $t = 2$, this is our new prior: $p(\mu_2|y_1; \theta) = N(\mu_{2|1}, \Sigma_{2|1})$.

Filtering and one-step ahead prediction

- For $t = 2, \dots, T$, we recursively repeat the two steps:

Filtering

$$p(\mu_t | y_1, \dots, y_t; \theta) = \frac{p(y_t | \mu_t; \theta)p(\mu_t | y_{1:t-1}; \theta)}{p(y_{1:t}; \theta)}$$

One-step ahead prediction

$$p(\mu_{t+1} | y_{1:t}; \theta) = \int p(\mu_{t+1} | \mu_t; \theta)p(\mu_t | y_{1:t}; \theta) d\mu_t$$

- The predictive distribution $p(\mu_{t+1} | y_{1:t}; \theta)$ is the prior at the next iteration, i.e. at time $t + 1$

The Kalman filter for the Local Level Model

- The Kalman filter recursively calculates these two steps
- Start with the initial conditions $\mu_{1|0}$ and $\Sigma_{1|0}$
- For $t = 1, \dots, T$

$$v_t = y_t - \mu_{t|t-1},$$

$$F_t = \Sigma_{t|t-1} + \sigma_\eta^2,$$

$$K_t = \Sigma_{t|t-1} / F_t,$$

$$\mu_{t|t} = \mu_{t|t-1} + K_t v_t, \quad \text{Filter step}$$

$$\Sigma_{t|t} = \Sigma_{t|t-1} - K_t \Sigma_{t|t-1}$$

$$\mu_{t+1|t} = \mu_{t|t}, \quad \text{Prediction step}$$

$$\Sigma_{t+1|t} = \Sigma_{t|t} + \sigma_\varepsilon^2.$$

Filtering and one-step ahead prediction

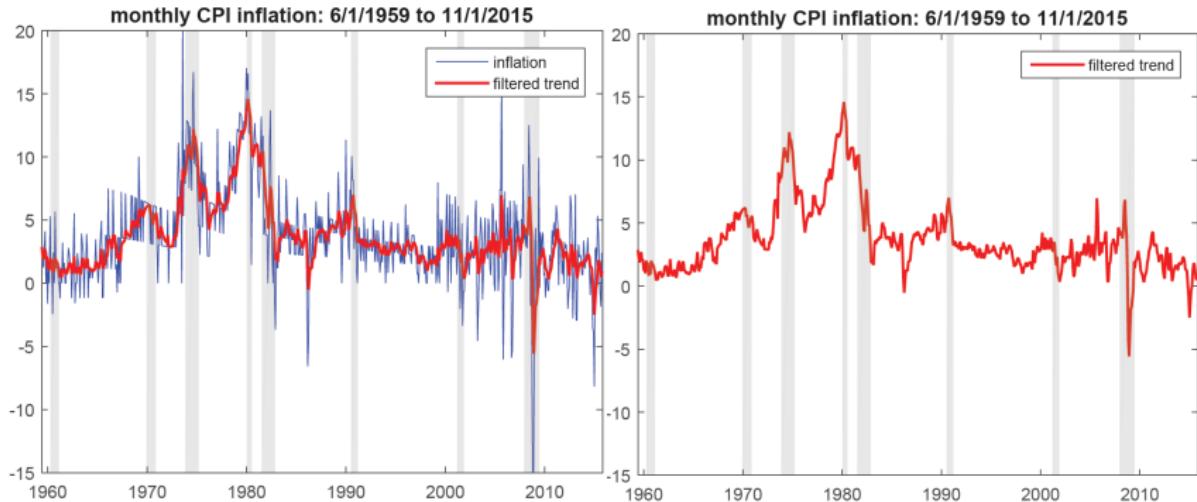
- These distributions describe our uncertainty about μ_t conditional on different information sets.
- All distributions are Gaussian!
- Only calculate means and covariance matrices (sufficient statistics)
- The filtering distribution

$$p(\mu_t | y_1, \dots, y_t; \theta) = N(\mu_{t|t}, \Sigma_{t|t}) \quad t = 1, \dots, T$$

- One-step ahead predictive distribution

$$p(\mu_{t+1} | y_1, \dots, y_t; \theta) = N(\mu_{t+1|t}, \Sigma_{t+1|t}) \quad t = 0, \dots, T$$

U.S. inflation



U.S. inflation, monthly continuously compounded. Left: inflation & filtered trend. Right: filtered trend.

The State Space Model (SSM)

Linear, Gaussian State Space Models

- Linear, Gaussian state space models are widely applicable
- For example, the following models can be placed in **state space form**
 - ① local level model
 - ② AR(1) observed in noise
 - ③ ARMA(p, q) models
 - ④ VAR(p) models
 - ⑤ Linear regression with serially correlated errors
 - ⑥ Time-varying parameter models
 - ⑦ Structural time series models
 - ⑧ Many more!
- First, we need to generalize the model.

Linear, Gaussian state space models

Definition

A **linear, Gaussian state space model** has observation equation

$$y_t = Z_t \alpha_t + d_t + \eta_t \quad \eta_t \sim N(0, H_t)$$

state transition equation

$$\alpha_{t+1} = T_t \alpha_t + c_t + R_t \varepsilon_{t+1} \quad \varepsilon_{t+1} \sim N(0, Q_t)$$

and initial conditions

$$\alpha_1 \sim N(a_{1|0}, P_{1|0}).$$

- The **state variable** or **state vector** is α_t .
- The **system matrices** $Z_t, d_t, H_t, T_t, c_t, R_t, Q_t$ are often time invariant: Z, d, H, T, c, R, Q .

Local level model

Consider the local level model

$$y_t = \mu_t + \eta_t, \quad \eta_t \sim N(0, \sigma_\eta^2)$$
$$\mu_{t+1} = \mu_t + \varepsilon_{t+1}, \quad \varepsilon_{t+1} \sim N(0, \sigma_\varepsilon^2)$$

The model can be written in state space form as:

$$\alpha_t = \mu_t, \quad Z = 1 \quad d = 0 \quad H = \sigma_\eta^2$$
$$T = 1 \quad c = 0 \quad R = 1 \quad Q = \sigma_\varepsilon^2$$

AR(1) observed in noise

$$\begin{aligned}y_t &= \mu_t + \eta_t & \eta_t &\sim N(0, \sigma_\eta^2) \\ \mu_{t+1} &= \phi_0 + \phi_1 \mu_t + \varepsilon_{t+1} & \varepsilon_{t+1} &\sim N(0, \sigma_\varepsilon^2)\end{aligned}$$

The model can be written in state space form as:

$$\alpha_t = \mu_t \quad Z = 1 \quad d = 0 \quad H = \sigma_\eta^2$$

$$T = \phi_1 \quad c = \phi_0 \quad R = 1 \quad Q = \sigma_\varepsilon^2$$

(Note: the local level model sets $\phi_0 = 0$ and $\phi_1 = 1$.)

ARMA(3,2) (version 1)

$$y_t = \mu + \phi_1(y_{t-1} - \mu) + \phi_2(y_{t-2} - \mu) + \phi_3(y_{t-3} - \mu) + \varepsilon_t \\ + \vartheta_1\varepsilon_{t-1} + \vartheta_2\varepsilon_{t-2} \quad \varepsilon_t \sim N(0, \sigma_\varepsilon^2)$$

Let $\phi_0 = (1 - \phi_1 - \phi_2 - \phi_3)\mu$. A state space form is:

$$Z = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 \end{pmatrix} \quad d = 0 \quad H = 0 \quad Q = \sigma_\varepsilon^2$$
$$\alpha_t = \begin{pmatrix} y_t \\ y_{t-1} \\ y_{t-2} \\ \varepsilon_t \\ \varepsilon_{t-1} \end{pmatrix} \quad T = \begin{pmatrix} \phi_1 & \phi_2 & \phi_3 & \vartheta_1 & \vartheta_2 \\ 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \end{pmatrix} \quad c = \begin{pmatrix} \phi_0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix} \quad R = \begin{pmatrix} 1 \\ 0 \\ 0 \\ 1 \\ 0 \end{pmatrix}$$

ARMA(3,2) (version 2)

$$y_t = \mu + \phi_1(y_{t-1} - \mu) + \phi_2(y_{t-2} - \mu) + \phi_3(y_{t-3} - \mu) + \varepsilon_t \\ + \vartheta_1\varepsilon_{t-1} + \vartheta_2\varepsilon_{t-2} \quad \varepsilon_t \sim N(0, \sigma_\varepsilon^2)$$

Let $\phi_0 = (1 - \phi_1 - \phi_2 - \phi_3)\mu$. An alternative state space form is:

$$Z = \begin{pmatrix} 1 & 0 & \dots & 0 \end{pmatrix} \quad d = 0 \quad H = 0 \quad Q = \sigma_\varepsilon^2$$

$$\alpha_t = \begin{pmatrix} y_t \\ \phi_2 y_{t-1} + \phi_3 y_{t-2} + \vartheta_1 \varepsilon_t + \vartheta_2 \varepsilon_{t-1} \\ \phi_3 y_{t-1} + \vartheta_2 \varepsilon_t \end{pmatrix}$$

$$T = \begin{pmatrix} \phi_1 & 1 & 0 \\ \phi_2 & 0 & 1 \\ \phi_3 & 0 & 0 \end{pmatrix} \quad c = \begin{pmatrix} \phi_0 \\ 0 \\ 0 \end{pmatrix} \quad R = \begin{pmatrix} 1 \\ \vartheta_1 \\ \vartheta_2 \end{pmatrix}$$

Note: the dimension of α_t is smaller than the last slide.

Remarks

- As the example of an ARMA(3, 2) shows us, the state space form of a model is not unique.
- There are multiple ways to place the same model in state space form. The definition of the 'state vector' is not necessarily the same in each case.
- The ARMA(3, 2) example also shows us that the state variable α_t is not always a latent variable.
- For examples on how to write ARMA(p, q) models in state space form; see Hamilton (1994) or Durbin and Koopman (2012).

VAR(p) in state-space form

$$y_t = \Phi_0 + \Phi_1 y_{t-1} + \dots + \Phi_p y_{t-p} + \varepsilon_t \quad \varepsilon_t \sim N(0, \Sigma_\varepsilon)$$

The model can be written in state space form as:

$$Z = \begin{pmatrix} I & 0 & \dots & 0 \end{pmatrix} \quad d = 0 \quad H = 0 \quad Q = \Sigma_\varepsilon$$
$$\alpha_t = \begin{pmatrix} y_t \\ y_{t-1} \\ \vdots \\ y_{t-p} \end{pmatrix}, \quad T = \begin{pmatrix} \Phi_1 & \Phi_2 & \dots & \Phi_p \\ I & 0 & \dots & 0 \\ 0 & I & \dots & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & I & 0 \end{pmatrix}, \quad c = \begin{pmatrix} \Phi_0 \\ 0 \\ \vdots \\ 0 \end{pmatrix}, \quad R = \begin{pmatrix} I \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

Linear regression with serially correlated errors

We can estimate optimally the β in a regression with AR(1) errors (you can use more complicated forms), as opposed to just accounting for it using HAC standard errors:

$$y_t = X_t \beta + \nu_t$$
$$\nu_{t+1} = \phi \nu_t + \varepsilon_{t+1} \quad \varepsilon_{t+1} \sim N(0, \sigma_\varepsilon^2)$$

The model can be written in state space form as:

$$\begin{aligned}\alpha_t &= \nu_t & Z &= 1 & d_t &= X_t \beta & H &= 0 \\ T &= \phi & c &= 0 & R &= 1 & Q &= \sigma_\varepsilon^2\end{aligned}$$

Time-varying parameter models

Here, the slopes are changing over time! Note, as in standard OLS, X_t must be observed at time t

$$\begin{aligned}y_t &= X_t \beta_t + \eta_t & \eta_t &\sim N(0, \Omega) \\ \beta_{t+1} &= \Phi_0 + \Phi_1 \beta_t + \varepsilon_{t+1} & \varepsilon_t &\sim N(0, \Sigma_\varepsilon)\end{aligned}$$

The model can be written in state space form as:

$$\begin{aligned}\alpha_t &= \beta_t & Z_t &= X_t & d &= 0 & H &= \Omega \\ T &= \Phi_1 & c &= \Phi_0 & R &= I & Q &= \Sigma_\varepsilon\end{aligned}$$

Note: a special case of this is a CAPM with time-varying β .

Missing data

Let's say that at time t some elements in

$$y_t = Z_t \alpha_t + d_t + \eta_t \quad \eta_t \sim N(0, H_t)$$

are missing.

- Simply set the variance of each missing observed component in H_t to infinity (in practice, a really big number), and zero its covariances. Alternatively, drop those rows of y_t , Z_t , d_t , and H_t in the update so that F_t stays nonsingular.

In Bayesian updating, if a signal has infinite variance of noise, the weight on that signal is zero.

- Next, when we see the updating equations, you will see H_t^{-1} arise which accounts for this effect

Predictive systems

We can use the SSM also on a cross-section for a big-data implementation

- ① Let $y_t = R_t^e$ be vector of $N \times 1$ of firms. Can even make this $N_t \times 1$, just make sure Z_t matrix reflects this at each t .
- ② Let the i 'th row of Z_t be $Z_t^{(i)} = [1 \ X'_{i,t-1}]$, where $X_{i,t-1}$ is a $K \times 1$ vector of characteristics

Now, there are $K + 1$ latent variables, where the loading on the last K depends on the associated characteristics of the firm at time t .

Be careful of specification of H_t and Q_t , especially for interpretation of estimated standard errors

- But, can view this as a quasi-maximum likelihood approach where we use out-of-sample performance instead of standard errors to evaluate model performance

Important things to remember

- There is more than one way to place a model in **state space form**
- Consequently, the definition of the '**state variable**' may change depending on how you do it.
- No matter how you place a model in **state space form** some things will not change:
 - ① forecasts of future data y_{t+h} for $h > 0$
 - ② the log-likelihood of the model

The Kalman filter calculates....

- The filtering distribution

$$p(\alpha_t | y_{1:t}; \theta) = \frac{p(y_t | \alpha_t; \theta)p(\alpha_t | y_{1:t-1}; \theta)}{p(y_{1:t}; \theta)}$$

- One-step ahead predictive distribution

$$p(\alpha_{t+1} | y_{1:t}; \theta) = \int p(\alpha_{t+1} | \alpha_t; \theta)p(\alpha_t | y_{1:t}; \theta) d\alpha_t$$

- All distributions are Gaussian! We only need their means and covariance matrices (their sufficient statistics).
- Our notation:

$$p(\alpha_t | y_{1:t}; \theta) = N(a_{t|t}, P_{t|t})$$

$$p(\alpha_{t+1} | y_{1:t}; \theta) = N(a_{t+1|t}, P_{t+1|t})$$

The Kalman Filter

- The Kalman filter recursively calculates these two steps.
- Start with the initial conditions $a_{1|0}$ and $P_{1|0}$.
- For $t = 1, \dots, T$

$$v_t = y_t - Z_t a_{t|t-1} - d_t,$$

$$F_t = Z_t P_{t|t-1} Z_t' + H_t,$$

$$K_t = P_{t|t-1} Z_t' F_t^{-1},$$

$$a_{t|t} = a_{t|t-1} + K_t v_t, \quad \text{Filter step}$$

$$P_{t|t} = P_{t|t-1} - K_t Z_t P_{t|t-1},$$

$$a_{t+1|t} = T_t a_{t|t} + c_t, \quad \text{Prediction step}$$

$$P_{t+1|t} = T_t P_{t|t} T_t' + R_t Q_t R_t'$$

Prediction and then filtering

- Some researchers reverse the order of the steps.
- Start with the initial (filtering) conditions $a_{0|0}$ and $P_{0|0}$.
- For $t = 1, \dots, T$

$$a_{t|t-1} = T_{t-1} a_{t-1|t-1} + c_{t-1}, \quad \text{Prediction step}$$

$$P_{t|t-1} = T_{t-1} P_{t-1|t-1} T'_{t-1} + R_{t-1} Q_{t-1} R'_{t-1}$$

$$v_t = y_t - Z_t a_{t|t-1} - d_t,$$

$$F_t = Z_t P_{t|t-1} Z'_t + H_t,$$

$$K_t = P_{t|t-1} Z'_t F_t^{-1},$$

$$a_{t|t} = a_{t|t-1} + K_t v_t, \quad \text{Filter step}$$

$$P_{t|t} = P_{t|t-1} - K_t Z_t P_{t|t-1},$$

The Kalman Predictor

- **Kalman predictor:** The filtered values $a_{t|t}$ and $P_{t|t}$ are never calculated.
- Start with the initial conditions $a_{1|0}$ and $P_{1|0}$.
- For $t = 1, \dots, T$

$$v_t = y_t - Z_t a_{t|t-1} - d_t,$$

$$F_t = Z_t P_{t|t-1} Z_t' + H_t,$$

$$M_t = T_t P_{t|t-1} Z_t' F_t^{-1},$$

$$L_t = T_t - M_t Z_t,$$

$$a_{t+1|t} = T_t a_{t|t-1} + c_t + M_t v_t,$$

$$P_{t+1|t} = T_t P_{t|t-1} L_t' + R_t Q_t R_t'$$

- Computationally faster because $a_{t|t}$ and $P_{t|t}$ are not calculated.

Initializing the Kalman filter

- We need values for $a_{1|0}$ and $P_{1|0}$ to start the Kalman filtering recursions.
- There are many suggestions in the literature for how to choose these values.
You should think of them as part of your model!
- In practice, we encounter two common situations:
 - ① α_t is stationary
 - ② α_t is non-stationary
- Case 1 is pretty easy. Case 2 is not.
- In the literature, $a_{1|0}$ and $P_{1|0}$ are often called **initial conditions**.

Stationarity of the state equation

- Suppose the state equation α_t is a stationary process.
- This means there exists a stationary (marginal) distribution $N(\mu_\alpha, V_\alpha)$.
- Stationarity means that all the eigenvalues of the matrix T are inside the unit circle.

Initializing the Kalman filter (stationary case)

- Suppose the state equation α_t is a stationary process.
- Let $N(\mu_\alpha, V_\alpha)$ denote the stationary distribution of α_t .
- Taking unconditional expectations, we find

$$E[\alpha_{t+1}] = TE[\alpha_t] + c + RE[\varepsilon_t]$$

$$\mu_\alpha = T\mu_\alpha + c$$

$$\Rightarrow \mu_\alpha = (I - T)^{-1}c$$

- If we use the stationary distribution, we set the mean to be

$$a_{1|0} = \mu_\alpha$$

Initializing the Kalman filter (stationary case)

- Let $N(\mu_\alpha, V_\alpha)$ denote the stationary distribution of α_t .
- Taking unconditional variances, we find

$$V[\alpha_{t+1}] = TV[\alpha_t]T' + RV[\varepsilon_t]R'$$

$$V_\alpha = TV_\alpha T' + RQR'$$

$$\text{vec}(V_\alpha) = \text{vec}(TV_\alpha T') + \text{vec}(RQR')$$

$$\text{vec}(V_\alpha) = (T \otimes T)\text{vec}(V_\alpha) + \text{vec}(RQR')$$

$$\Rightarrow \text{vec}(V_\alpha) = [I - (T \otimes T)]^{-1} \text{vec}(RQR')$$

- If we use the stationary distribution, we set the covariance matrix as

$$P_{1|0} = V_\alpha$$

Initializing the Kalman filter (non-stationary)

- In some models of interest, elements of the state vector α_t are non-stationary.
- These models have unit roots; e.g. the local level model.
- The stationary distribution of α_t does not exist.
- The simple way to initialize the Kalman filter is set the variance $P_{1|0}$ to a really large number; e.g. 10^4
- Koopman (1997) gives an exact initialization. Effectively, it calculates a conditional log-likelihood function that drops the initial parts of the likelihood. Not easy to implement though. See also Chapter 5 of Durbin and Koopman (2012).

Impact of the initial conditions

- If the model is stationary, the **initial conditions** $a_{1|0}$ and $P_{1|0}$ typically do not have a large influence on the results.
- Filtered estimates $a_{t|t}$ and $P_{t|t}$ will converge to the same thing (they are equal) even if we start the Kalman filter from different initial conditions!!
- This is due to the stationarity of the model.
- The early estimates $a_{t|t}$ and $P_{t|t}$ will be different during the first few iterations: $t < 20$ or so.
- For non-stationary models, the initial conditions can have an impact if the overall (time-series) sample size T is small.

Smoothing distributions

- During the forward pass of the Kalman filter we calculate the filtering and one-step ahead predictive distributions.

$$p(\alpha_t | y_{1:t}; \theta) = N(a_{t|t}, P_{t|t})$$

$$p(\alpha_{t+1} | y_{1:t}; \theta) = N(a_{t+1|t}, P_{t+1|t})$$

for $t = 1, \dots, T$

- These distributions describe our uncertainty about the state α_t conditional on different information sets.
- For many time series models, there is still uncertainty about the state vector α_t even after we observe **all** the data.

$$p(\alpha_t | y_1, \dots, y_T; \theta)$$

- This is called the **smoothed distribution**.

Smoothing distributions

- At time $t = T$, the filtered and smoothed distributions are equal!

$$p(\alpha_T | y_{1:T}; \theta) = N(a_{T|T}, P_{T|T})$$

- At $t = T$, we know the mean $a_{T|T}$ and covariance matrix $P_{T|T}$.
- We can write the sufficient statistics $a_{t|T}, P_{t|T}$ as a recursive function of $a_{t+1|T}, P_{t+1|T}$.
- The **Kalman smoother** recursively calculates the smoothing distributions backwards

$$p(\alpha_t | y_{1:T}; \theta) = N(a_{t|T}, P_{t|T})$$

for $t = T - 1, \dots, 1$

Kalman smoother (Rauch-Tung-Striebel)

- Run the Kalman filter forward in time for $t = 1, \dots, T$.
- Store the quantities $\left\{ a_{t|t-1}, P_{t|t-1}, v_t, F_t, L_t \right\}_{t=1}^T$.
- Set $r_{T+1} = 0$ and $N_{T+1} = 0$
- For $t = T, \dots, 1$

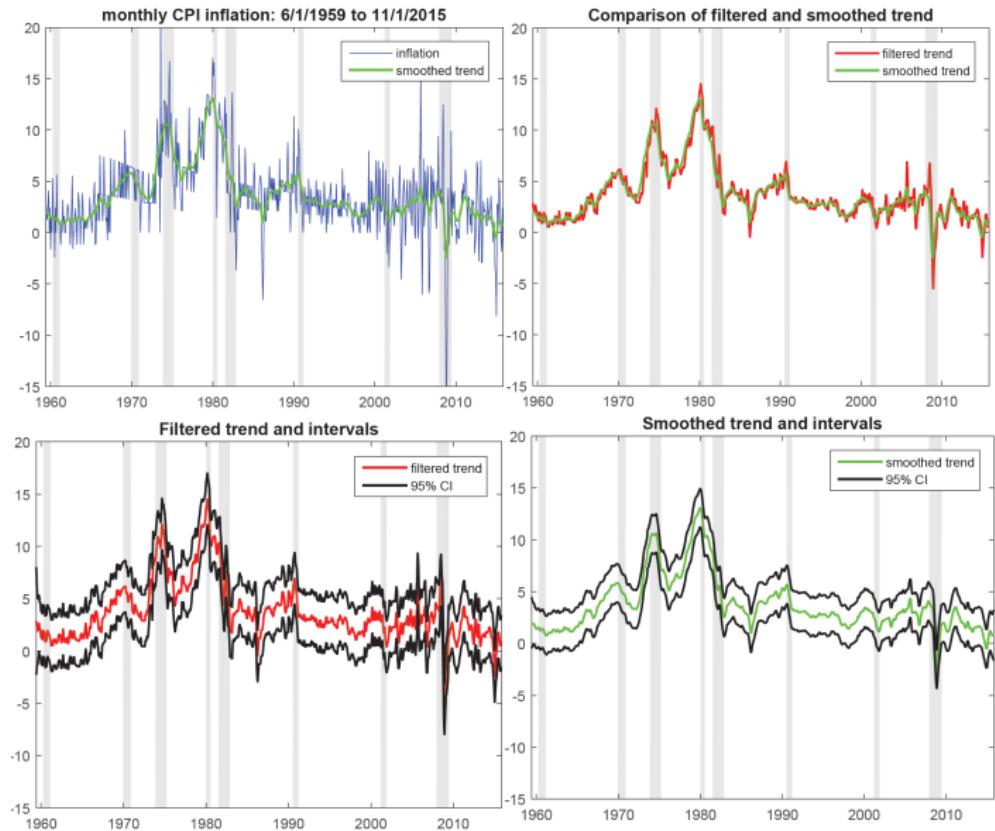
$$r_t = Z_t' F_t^{-1} v_t + L_t' r_{t+1}$$

$$N_t = Z_t' F_t^{-1} Z_t + L_t' N_{t+1} L_t$$

$$a_{t|T} = a_{t|t-1} + P_{t|t-1} r_t$$

$$P_{t|T} = P_{t|t-1} - P_{t|t-1} N_t P_{t|t-1}$$

Example: U. S. inflation



U.S. inflation, monthly continuously compounded.

Comments on the Kalman smoother

- There are different versions of the **Kalman smoother**.
- All of them compute the means $a_{t|T}$ and covariance matrices $P_{t|T}$.
- They differ depending on what you store in computer memory on the forward pass of the Kalman filter.
- In the statistics/econometrics/engineering literature, the terms '**filtering**' and '**smoothing**' are used in different ways.
 - ▶ in state space models, filtering means to use data only up to time t .
 - ▶ in state space models, smoothing means to use ALL the data T .
 - ▶ however, different parts of science use the terms '**filtering**' and '**smoothing**' to mean other things.

Forecasting

- We often want to forecast:
 - ➊ the state variable h steps ahead: α_{t+h}
 - ➋ future data h steps ahead: y_{t+h}
- We also want to characterize our uncertainty of these variables.
- This is easy to do in a state space model.

Forecasting the state variable

- To forecast the state variable α_t at time $t + h$, we need to calculate the predictive distribution.

$$p(\alpha_{t+h} | y_{1:t}; \theta)$$

- Under the assumption that the errors η_t and ε_t are Gaussian, the predictive distribution is Gaussian.

$$p(\alpha_{t+h} | y_{1:t}; \theta) = N(a_{t+h|t}, P_{t+h|t})$$

- We can calculate the mean $a_{t+h|t}$ and covariance matrix $P_{t+h|t}$.
- These can be calculated recursively: $t + 1$, then $t + 2$, then... $t + h$

Forecasting the state variable

- Let's assume that the model is time-invariant. The **system matrices** are constant

$$Z, d, H, T, c, R, Q$$

- At the end of the Kalman filter, we already have

$$p(\alpha_{t+1}|y_{1:t}; \theta) = N(a_{t+1|t}, P_{t+1|t})$$

- The next predictive distribution is:

$$p(\alpha_{t+2}|y_{1:t}; \theta) = \int p(\alpha_{t+2}|\alpha_{t+1}; \theta)p(\alpha_{t+1}|y_{1:t}; \theta)d\alpha_{t+1}$$

- To calculate $a_{t+2|t}, P_{t+2|t}$, we apply the recursion

$$a_{t+2|t} = Ta_{t+1|t} + c,$$

$$P_{t+2|t} = TP_{t+1|t}T' + RQR'$$

Forecasting the state variable

- To calculate these quantities at longer horizons, we simply iterate

$$a_{t+h|t} = Ta_{t+h-1|t} + c,$$

$$P_{t+h|t} = TP_{t+h-1|t}T' + RQR'$$

- If the system matrices are time-varying, you need to know their future values:

$$H_{t+j}, T_{t+j}, c_{t+j}, R_{t+j}, Q_{t+j} \quad j = 0, \dots, h-1$$

Forecasting future data

- To forecast future data y_{t+h} , do the following:
 - calculate the predictive distribution of α_{t+h} using the earlier recursion

$$\alpha_{t+h|t} = T\alpha_{t+h-1|t} + c,$$

$$P_{t+h|t} = TP_{t+h-1|t}T' + RQR'$$

- calculate the predictive mean and variance of y_{t+h} via

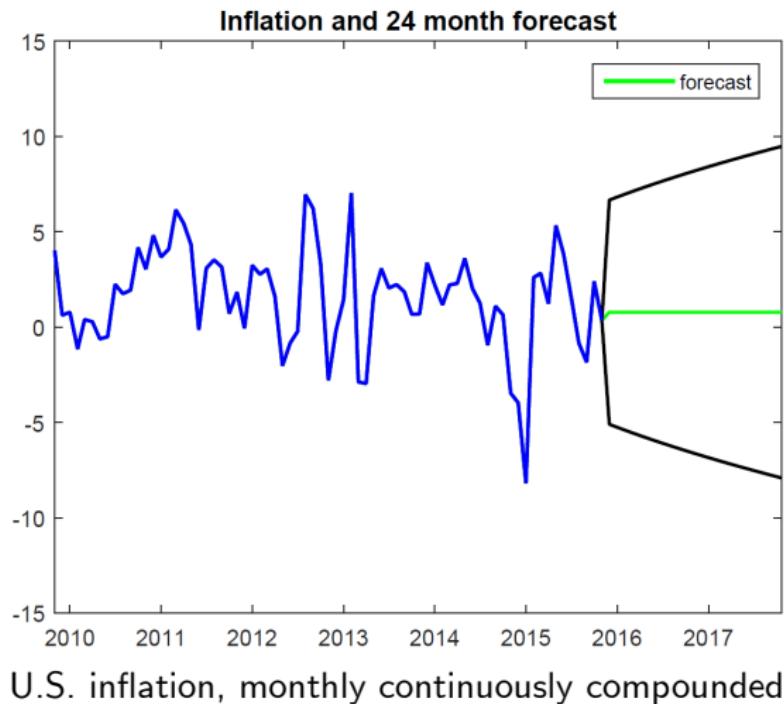
$$\hat{y}_{t+h|t} = Z\alpha_{t+h|t} + d$$

$$F_{t+h|t} = ZP_{t+h|t}Z' + H$$

- If the system matrices are time-varying, you need to know their future values:

$$Z_{t+j}, d_{t+j}, H_{t+j}, T_{t+j}, c_{t+j}, R_{t+j}, Q_{t+j} \quad j = 0, \dots, h-1$$

Example: U.S. inflation and the local level model



Example: U.S. inflation: ARMA(1,1)

Suppose we consider an ARMA(1,1) for inflation

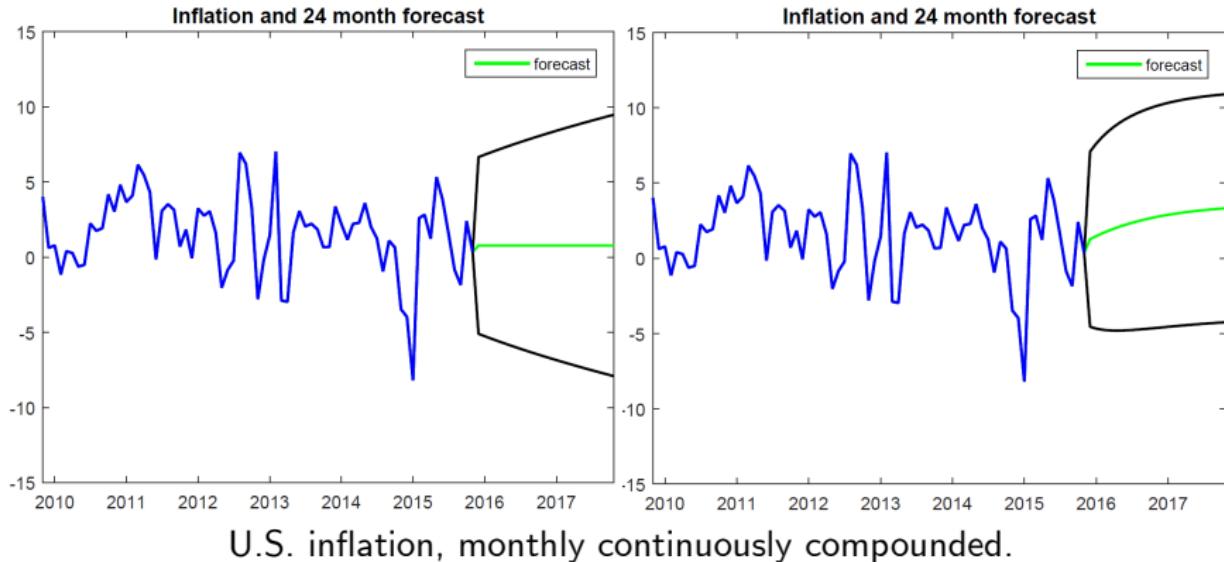
$$\pi_t = \mu + \phi_1(\pi_{t-1} - \mu) + \varepsilon_t + \theta_1 \varepsilon_{t-1} \quad \varepsilon_t \sim N(0, \sigma_\varepsilon^2)$$

I will use the state space form:

$$Z = \begin{pmatrix} 1 & 0 \end{pmatrix} \quad d = \mu \quad H = 0 \quad Q = \sigma_\varepsilon^2$$

$$\alpha_t = \begin{pmatrix} \pi_t - \mu \\ \varepsilon_t \end{pmatrix} \quad T = \begin{pmatrix} \phi_1 & \theta_1 \\ 0 & 0 \end{pmatrix} \quad c = \begin{pmatrix} 0 \\ 0 \end{pmatrix} \quad R = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$$

Example: U.S. inflation



Parameter estimation

- Thus far, we have assumed that we know the parameters θ .
- The log-likelihood of the model is the joint distribution of the data:

$$\ln p(y_1, y_2, \dots, y_T | \theta) = \sum_{t=2}^T \ln p(y_t | y_{1:t-1}; \theta) + \ln p(y_1; \theta)$$

- The Kalman filter calculates the time t contribution to the log-likelihood

$$\ln p(y_t | y_{1:t-1}; \theta)$$

at each iteration during the forward pass.

- We can maximize the log-likelihood (numerically) to estimate the parameters θ .

What is the likelihood?

- The likelihood at time t is just our predicted value of y_t given information up to time $t - 1$

$$p(y_t | y_1, \dots, y_{t-1}; \theta)$$

- Under the assumption that the errors are Gaussian, this distribution is Gaussian. Calculate the mean and covariance matrix.
- In the last section, we just showed how to calculate the forecast of y_{t+h} given information up to time t !!!
- We need to calculate the forecast of y_t given information up to time $t - 1$.

How to calculate the log-likelihood

- Start with the initial conditions $a_{1|0}$ and $P_{1|0}$
- Initialize the log-likelihood: $\ell_0 = 0$
- For $t = 1, \dots, T$

$$v_t = y_t - Z_t a_{t|t-1} - d_t,$$

$$F_t = Z_t P_{t|t-1} Z_t' + H_t,$$

$$M_t = T_t P_{t|t-1} Z_t' F_t^{-1},$$

$$L_t = T_t - M_t Z_t,$$

$$a_{t+1|t} = T_t a_{t|t-1} + c_t + M_t v_t,$$

$$P_{t+1|t} = T_t P_{t|t-1} L_t' + R_t Q_t R_t'$$

$$\ell_t = \ell_{t-1} - \frac{N}{2} \log(2\pi) - \frac{1}{2} \log |F_t| - \frac{1}{2} v_t' F_t^{-1} v_t$$

Appendix: Updating from noisy signal – the details

Bayes Rule and Updates in Beliefs

Let's do some math!

- Start with two known distributions:

$$x \sim N(\mu_X, \sigma_X^2)$$
$$y|x \sim N(x, \sigma_{Y|X}^2)$$

- Here x corresponds to x_t and y to z_t in our slides at the beginning of the lecture

Bayes Rule and Updates in Beliefs

- We want to get to the distribution of $x|y$, so let's first multiply these two pdf's:

$$\begin{aligned} & \frac{1}{\sqrt{2\pi\sigma_X^2}} \exp\left\{ \frac{(x - \mu_X)^2}{2\sigma_X^2} \right\} \frac{1}{\sqrt{2\pi\sigma_{Y|X}^2}} \exp\left\{ \frac{(y - x)^2}{2\sigma_{Y|X}^2} \right\} \\ &= \frac{1}{2\pi\sqrt{\sigma_X^2\sigma_{Y|X}^2}} \exp\left\{ \frac{(x - \mu_X)^2}{2\sigma_X^2} + \frac{(y - x)^2}{2\sigma_{Y|X}^2} \right\} \\ &= \frac{1}{2\pi\sqrt{\sigma_X^2\sigma_{Y|X}^2}} \exp\left\{ \begin{array}{l} \frac{(x - \mu_X)^2}{2\sigma_X^2} \frac{(\sigma_X^{-2} + \sigma_{Y|X}^{-2})^{-1}/\sigma_X^2}{(\sigma_X^{-2} + \sigma_{Y|X}^{-2})^{-1}/\sigma_X^2} \dots \\ + \frac{(y - x)^2}{2\sigma_{Y|X}^2} \frac{(\sigma_X^{-2} + \sigma_{Y|X}^{-2})^{-1}/\sigma_{Y|X}^2}{(\sigma_X^{-2} + \sigma_{Y|X}^{-2})^{-1}/\sigma_{Y|X}^2} \end{array} \right\} \end{aligned}$$

Oh yeah... Algebra!

Continuing...

Define $k \equiv (\sigma_X^{-2} + \sigma_{Y|X}^{-2})^{-1}$:

$$\begin{aligned}& \frac{1}{2\pi\sqrt{\sigma_X^2\sigma_{Y|X}^2}} \exp \left\{ \frac{(x - \mu_X)^2}{2\sigma_X^2} \frac{k/\sigma_X^2}{k/\sigma_X^2} + \frac{(y - x)^2}{2\sigma_{Y|X}^2} \frac{k/\sigma_{Y|X}^2}{k/\sigma_{Y|X}^2} \right\} \\&= \frac{1}{2\pi\sqrt{\sigma_X^2\sigma_{Y|X}^2}} \exp \left\{ \frac{(x^2 - 2x\mu_X + \mu_X^2) k/\sigma_X^2 + (y^2 - 2yx + x^2) k/\sigma_{Y|X}^2}{2k} \right\} \\&= \frac{1}{2\pi\sqrt{\sigma_X^2\sigma_{Y|X}^2}} \exp \left\{ \frac{x^2 k/\sigma_X^2 - 2x\mu_X k/\sigma_X^2 + \mu_X^2 k/\sigma_X^2 + y^2 k/\sigma_{Y|X}^2 - 2yxk/\sigma_{Y|X}^2 + x^2 k/\sigma_{Y|X}^2}{2k} \right\} \\&= \frac{1}{2\pi\sqrt{\sigma_X^2\sigma_{Y|X}^2}} \exp \left\{ \frac{x^2 k (\sigma_X^{-2} + \sigma_{Y|X}^{-2}) - 2x (yk/\sigma_{Y|X}^2 + \mu_X k/\sigma_X^2) + \mu_X^2 k/\sigma_X^2 + y^2 k/\sigma_{Y|X}^2}{2k} \right\}\end{aligned}$$

Finger-lickin'!

Continuing... ...

Note that $k \left(\sigma_X^{-2} + \sigma_{Y|X}^{-2} \right) = 1$. So:

$$\frac{1}{2\pi\sqrt{\sigma_X^2\sigma_{Y|X}^2}} \exp \left\{ \frac{x^2 - 2x \left(yk/\sigma_{Y|X}^2 + \mu_X k/\sigma_X^2 \right) + \mu_X^2 k/\sigma_X^2 + y^2 k/\sigma_{Y|X}^2}{2k} \right\}.$$

Next, complete the square:

$$\begin{aligned} & \frac{1}{\sqrt{2\pi k}} \exp \left\{ \frac{\left(x - \left(yk/\sigma_{Y|X}^2 + \mu_X k/\sigma_X^2 \right) \right)^2}{2k} \right\} \\ & \times \frac{\sqrt{2\pi k}}{2\pi\sqrt{\sigma_X^2\sigma_{Y|X}^2}} \exp \left\{ \frac{-\left(yk/\sigma_{Y|X}^2 + \mu_X k/\sigma_X^2 \right)^2 + \mu_X^2 k/\sigma_X^2 + y^2 k/\sigma_{Y|X}^2}{2k} \right\}. \end{aligned}$$

Note that the first line says $x|y$ is normally distributed with mean $\left(yk/\sigma_{Y|X}^2 + \mu_X k/\sigma_X^2 \right)$ and variance k .

- The second line is a constant (not a function of x), conditional on y . Since we only were given the distribution up to a proportion (recall the Bayes Rule equation), we can ignore it for our purposes.

Learning with Normal Distributions

In sum, we are looking for the distribution of x conditional on a data point, y .

We found that $x|y$ is normally distributed using Bayes Rule.

- The mean of this distribution is:

$$\begin{aligned} yk/\sigma_{Y|X}^2 + \mu_X k/\sigma_X^2 &= y \frac{\sigma_{Y|X}^{-2}}{\sigma_X^{-2} + \sigma_{Y|X}^{-2}} + \mu_X \frac{\sigma_X^{-2}}{\sigma_X^{-2} + \sigma_{Y|X}^{-2}} \\ &= y \times (1 - \text{weight on prior}) + \mu_X \times (\text{weight on prior}) \end{aligned}$$

Note that the more precise the signal is (the higher $\sigma_{Y|X}^{-2}$ is) and the less precise the prior is (the lower σ_X^{-2} is), the more weight is given to the signal when updating the mean belief about x .

- The variance is $k = \left(\sigma_X^{-2} + \sigma_{Y|X}^{-2}\right)^{-1} < \sigma_X^2$.

References

- Burmeister, E., K. D. Wall, and J. D. Hamilton (1986). Estimation of unobserved expected monthly inflation using Kalman Filtering. *Journal of Business and Economic Statistics* 4(2), 147-160.
- Durbin, J. and S. J. Koopman (2012). Time Series Analysis by State Space Methods (2 ed.). Oxford, UK: Oxford University Press.
- Hamilton, J. D. (1994). Time Series Analysis. Princeton, NJ: Princeton University Press.
- Kalman, R. (1960). A new approach to linear Itering and prediction problems. *Journal of Basic Engineering, Transactions of the ASME* 82(1), 35-46.
- Koopman, S. J. (1997). Exact initial Kalman Filtering and smoothing for non-stationary time series models. *Journal of the American Statistical Association* 92(440), 1630-1638.
- Stock, J. H. and M. W. Watson (2007). Why has U.S. inflation become harder to forecast? *Journal of Money, Credit, and Banking* 39(1), 3-34.