

Lecture 2

Thinking in “Regimes” in economics and finance

Hidden Markov Models

Lars A. Lochstoer

UCLA Anderson School of Management

Overview

- Hidden Markov chain models
- Switching regime models
- Markov-switching State Space Models
- The Unscented Kalman filter for nonlinear SSMs

Motivation

Many economic phenomena can be described as “regimes”

- ① Recessions and expansions
- ② Bull/bear markets
- ③ Financial crisis
- ④ Currency regimes (fixed, floating, crawling peg)
- ⑤ High inflation, low inflation
- ⑥ Uncertainty: technological, geopolitical, financial system, etc.

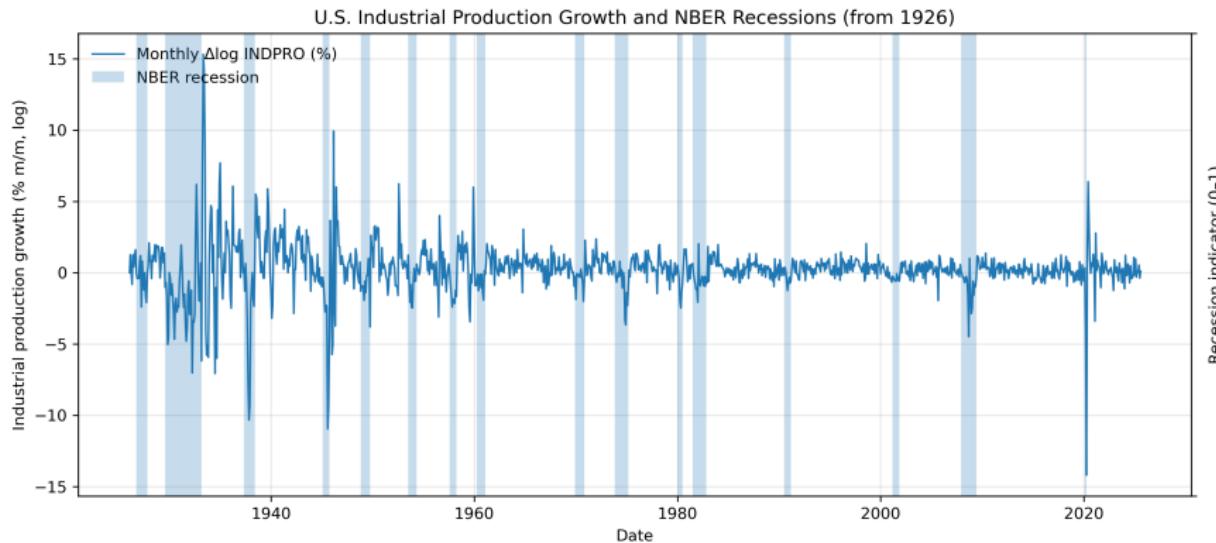
Common to these: nonlinear dynamics

- Large changes vs small changes
- Different expected duration of different regimes

Motivation: The business cycle

Industrial production data 1926-2025

- Think of this as economic growth data
- Shaded areas denote recessions as determined by the National Bureau of Economic Research (NBER)



Hidden Markov Models

The Switching Regime (SR) Model

Hidden Markov chain Models (HMMs)

A convenient tool: thinking in terms of states and state transitions

Latent chain (states) $S_t \in \{1, \dots, K\}$ with transition matrix

$$\Pi = \begin{bmatrix} \pi_{11} & \pi_{12} & \cdots & \pi_{1K} \\ \pi_{21} & \pi_{22} & \cdots & \pi_{2K} \\ \vdots & \vdots & \ddots & \vdots \\ \pi_{K1} & \pi_{K2} & \cdots & \pi_{KK} \end{bmatrix}$$

where $\Pr(S_t = j | S_{t-1} = i) = \pi_{ij}$ and $\sum_j \pi_{ij} = 1$.

- Observations y_t have dynamics governed by the state: pdf is $f_{S_t}(y_t; \theta_{S_t})$
- Need prior over the initial state, p_0
- Parameter set $\Theta = \{\Pi, p_0, \theta_1, \dots, \theta_K\}$.

Simplest case: Two-State Gaussian HMM

Let $f_{S_t}(y_t; \theta_{S_t})$ be $N(\mu_{s_t}, \sigma_{s_t}^2)$. That is:

$$y_t = \mu_{s_t} + \sigma_{s_t} \varepsilon_t,$$

where $\varepsilon_t \sim N(0, 1)$, and hidden regime $S_t \in \{1, 2\}$ with transition matrix

$$\Pi = \begin{bmatrix} \pi_{11} & 1 - \pi_{11} \\ 1 - \pi_{22} & \pi_{22} \end{bmatrix}$$

Thus, we have 6 parameters to estimate $(\pi_{11}, \pi_{22}, \mu_1, \mu_2, \sigma_1, \sigma_2)$, plus we need a prior for initial state, p_0 .

Forecasting with two-state Gaussian HMM

Current beliefs about the state: $p_t(s_t)$ (could write, $p_{t|t}(s_t)$)

- E.g., $p_t(s_t = 1) = 0.7, p_t(s_t = 2) = 0.3$

Then

$$\begin{aligned}E_t(y_{t+1}) &= E_t(E_t(y_{t+1}|s_t)) \\&= p_t(s_t = 1) E_t(y_{t+1}|s_t = 1) + p_t(s_t = 2) E_t(y_{t+1}|s_t = 2)\end{aligned}$$

Intermediate calculation:

$$\begin{aligned}E_t(y_{t+1}|s_t = 1) &= \pi_{11}\mu_1 + (1 - \pi_{11})\mu_2 = e'_1 \Pi \mu, \\E_t(y_{t+1}|s_t = 2) &= (1 - \pi_{22})\mu_1 + \pi_{22}\mu_2 = e'_2 \Pi \mu,\end{aligned}$$

where $\mu = [\mu_1 \quad \mu_2]'$ and e_i is a column vector with a one in element i and zeros otherwise.

Forecasting (cont'd)

Repeating from last slide and define $\mathbf{p}_t = [p_t(s_t = 1) \quad p_t(s_t = 2)]'$.

$$\begin{aligned} E_t(y_{t+1}) &= p_t(s_t = 1) E_t(y_{t+1}|s_t = 1) + p_t(s_t = 2) E_t(y_{t+1}|s_t = 2) \\ &= \mathbf{p}'_t \Pi \mu, \end{aligned}$$

Longer-horizon forecasts:

$$E_t(y_{t+2}) = E_t(E_{t+1}(y_{t+2})) = E_t(\mathbf{p}'_{t+1} \Pi \mu)$$

Thus, we have to forecast future beliefs, \mathbf{p}_{t+1} :

$$E_t(\mathbf{p}_{t+1}) = \mathbf{p}_{t+1|t} = \Pi' \mathbf{p}_t$$

and so

$$E_t(y_{t+2}) = \mathbf{p}'_t \Pi^2 \mu$$

which generalizes to

$$E_t(y_{t+j}) = \mathbf{p}'_t \Pi^j \mu, \quad \text{where} \quad j = 1, 2, 3, \dots$$

Duration of regimes

Expected duration of regime i :

$$D_i = \min \{k \geq 1 : S_{t+k} \neq i | S_t = i\},$$

$$\Pr(D_i = k) = \pi_{ii}^{k-1} (1 - \pi_{ii}),$$

$$E(D_i) = \frac{1}{1 - \pi_{ii}}, \quad \text{Var}(D_i) = \frac{\pi_{ii}}{(1 - \pi_{ii})^2}$$

Expected length one stays in regime i is exploding as π_{ii} approaches 1

- Thus, with $\pi_{11} \neq \pi_{22}$ one can, as in the data, have recessions (typically) last for a shorter time than expansions
- This is a form of nonlinearity. Compare to AR(1) for instance.

Model estimation

The steps to estimate an HMM share many similarities with the Kalman filter

- Smoothing and filtering probabilities
- Bayes rule
- But, this time not a linear Gaussian setup, which causes differences

We will discuss these steps in order

- ① Hamilton filter for forward filtering probabilities
- ② Likelihood function
- ③ Smoothing probabilities and the expectation (E) step
- ④ Parameter estimates and the maximization (M) step

Hamilton Filter: Forward Probabilities

Predictive regime probabilities:

$$p_{t|t-1}(j) = \sum_{i=1}^2 \pi_{ij} p_{t-1}(i).$$

Observation density by regime:

$$f_j(y_t) = N(y_t; \mu_j, \sigma_j^2).$$

Marginal predictive density:

$$p(y_t | y_{1:t-1}) = \sum_{j=1}^2 f_j(y_t) p_{t|t-1}(j).$$

Bayes update:

$$p_t(j) = \frac{f_j(y_t) p_{t|t-1}(j)}{\sum_{m=1}^2 f_m(y_t) p_{t|t-1}(m)}.$$

Exact Log-Likelihood

The log-likelihood function is given by

$$\mathcal{L}(\Theta) = \sum_{t=1}^T \log \left(\sum_{j=1}^2 N(y_t; \mu_j, \sigma_j^2) p_{t|t-1}(j) \right).$$

Notice that we now at each time t have a probability weighted average of two Normals, corresponding to one for each regime.

EM Algorithm: E-Step

First, we get the smoothing probabilities given a guess of the parameters

$$p_{t|T}(j) = \Pr(S_t = j | y_{1:T}), \quad \xi_{t|T}(i,j) = \Pr(S_{t-1} = i, S_t = j | y_{1:T}).$$

Define the backward variable

$$\beta_t(j) = \Pr(y_{t+1}, y_{t+2}, \dots, y_T | S_t = j).$$

Then

$$p_{t|T}(j) \propto p_t(j) \beta_t(j), \quad \xi_{t|T}(i,j) \propto p_{t-1}(i) \pi_{ij} N(y_t; \mu_j, \sigma_j^2) \beta_t(j).$$

Normalize per t so probabilities sum to 1.

EM Algorithm: M-Step

$$\hat{\mu}_j = \frac{\sum_{t=1}^T y_t p_{t|T}(j)}{\sum_{t=1}^T p_{t|T}(j)},$$

$$\hat{\sigma}_j^2 = \frac{\sum_{t=1}^T (y_t - \hat{\mu}_j)^2 p_{t|T}(j)}{\sum_{t=1}^T p_{t|T}(j)},$$

$$\hat{\pi}_{ij} = \frac{\sum_{t=2}^T \xi_t(i, j)}{\sum_{t=2}^T p_{t-1|T}(i)}.$$

Notice that we take the simple mean and variance using the probability weights for each observation being in regime j .

- The transition probabilities simply count transitions, again using the relevant probability weights.

Iterated on the above 3 equations, until the parameter guesses (which enter into calculations on the right hand side of the equalities) equal the estimates on the left hand side.

- The EM algorithm is shown to maximize the likelihood function given earlier

Hints for MLE and Identification

- Maximize $\mathcal{L}(\Theta)$ with L-BFGS-B or Newton.
- Reparametrize: $\sigma_j^2 = \exp(\lambda_j)$, $\pi_{12} = \sigma(\alpha)$, $\pi_{21} = \sigma(\beta)$, where $\sigma(\cdot)$ refers to the Sigmoid function
- Label switching: impose $\mu_1 < \mu_2$ or $\sigma_1^2 < \sigma_2^2$.
 - ▶ Otherwise, there is not a unique maximum
 - ▶ You could label state 1 as state 2 and state 2 as state 1 and the model would yield the same fit

Forecasting: Summary

$$p_{t+1|t}(j) = \sum_{i=1}^2 \pi_{ij} p_t(i),$$

$$E(y_{t+1}|y_{1:t}) = \sum_{j=1}^2 \mu_j p_{t+1|t}(j),$$

$$V(y_{t+1}|y_{1:t}) = \sum_{j=1}^2 (\sigma_j^2 + \mu_j^2) p_{t+1|t}(j) - \left(\sum_{j=1}^2 \mu_j p_{t+1|t}(j) \right)^2.$$

Note that state persistence, and likelihood of being transitioned into, are important for how the state mean μ_j is reflected in future forecast

HMMs in Finance: Real-World Examples

Let's look at some HMMs from finance

Will stick to two regimes with iid dynamics conditional on regime

- ① Credit spread for high yield bonds
- ② Inflation dynamics

Credit: Low vs Stressed Spread Regimes

y_t is the credit spread level. $y_t \mid S_t = j \sim N(\mu_j, \sigma_j^2)$.

- Monthly data, estimate using ML as above

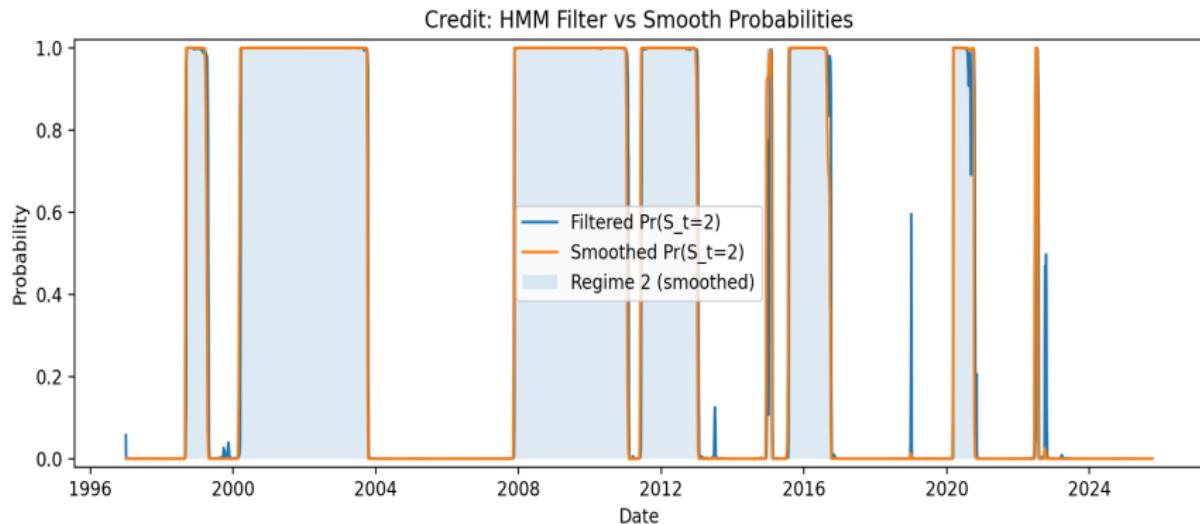
	μ_1	σ_1	μ_2	σ_2	π_{11}	π_{22}
<i>MLE</i>	3.81	0.71	7.57	2.72	0.998	0.997
(s.e.)	(0.013)	(0.009)	(0.055)	(0.036)	(0.001)	(0.001)

- Regime 1: low mean, low volatility
- Regime 2: high mean, high volatility
- Both regimes very persistent

Credit: Filtering and smoothing probabilities

y_t is the credit spread level. $y_t \mid S_t = j \sim N(\mu_j, \sigma_j^2)$.

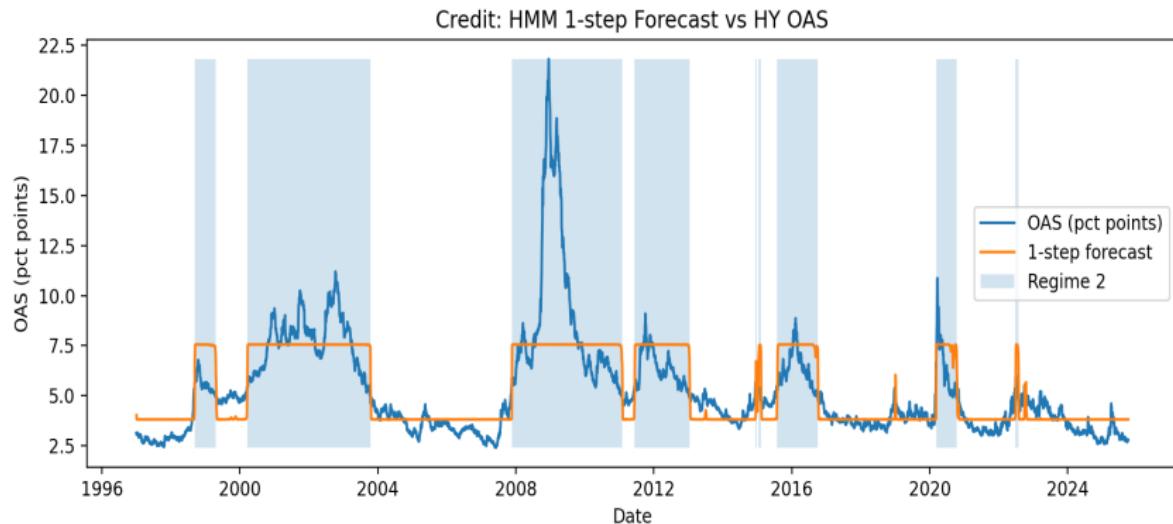
- Regime 1: low mean and vol
- Regime 2: high mean and vol
- Shaded area if $p_{t|T}(s_t = 2) > 0.5$



Credit: Low vs Stressed Spread Regimes

y_t is the credit spread level. $y_t \mid S_t = j \sim N(\mu_j, \sigma_j^2)$.

- Regime 1: low mean and vol
- Regime 2: high mean and vol
- Shaded area if $p_{t|T}(s_t = 2) > 0.5$



Inflation: Low vs High Volatility Regimes

- YoY inflation, but we will still use the same 2-state SR model here
- y_t is YoY monthly inflation. $y_t | S_t = j \sim N(\mu_j, \sigma_j^2)$.

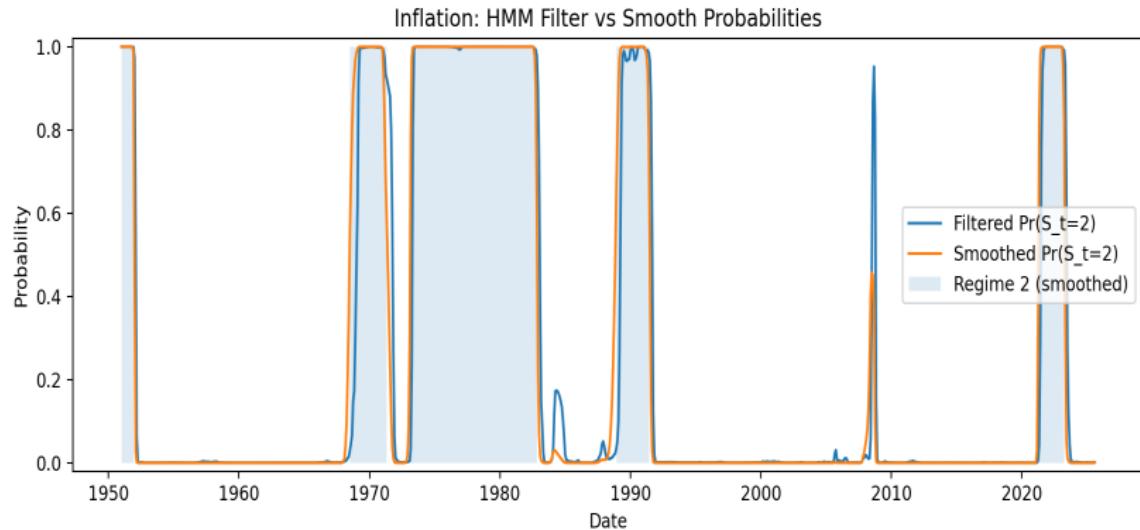
	μ_1	σ_1	μ_2	σ_2	π_{11}	π_{22}
MLE	2.27	1.21	7.44	2.65	0.992	0.978
(s.e.)	(0.055)	(0.038)	(0.215)	(0.127)	(0.004)	(0.010)

- Regime 1: low mean, low volatility
- Regime 2: high mean, high volatility
- Both regimes very persistent, but regime 2 less so
 - ▶ Regime 1: Duration 10.6 years; Regime 2: Duration 3.9 years.

Inflation: Filtering and smoothing probabilities

y_t is inflation. $y_t | S_t = j \sim N(\mu_j, \sigma_j^2)$.

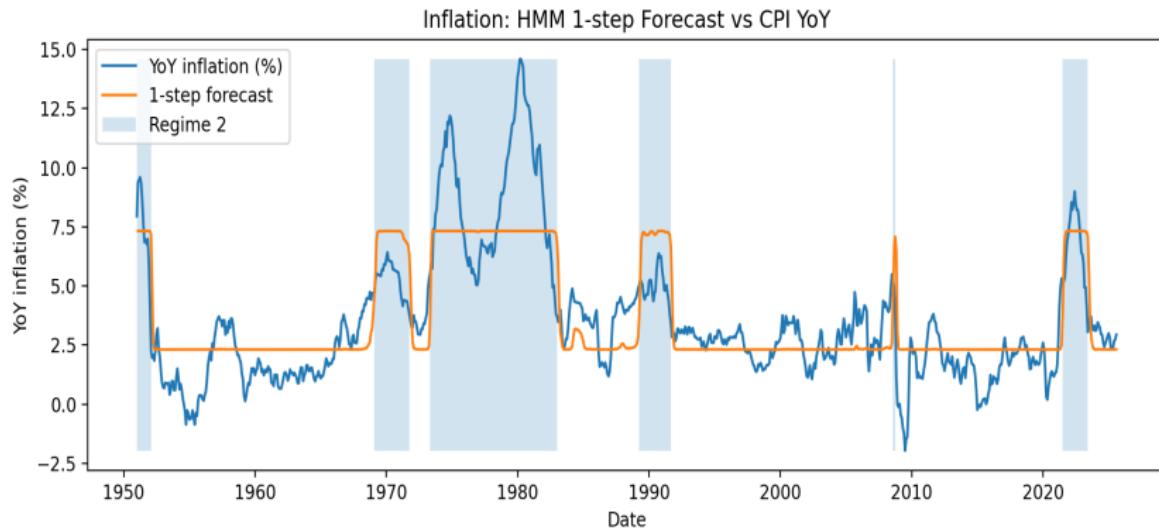
- Regime 1: low mean and vol
- Regime 2: high mean and vol, less persistent
- Shaded area if $p_{t|T}(s_t = 2) > 0.5$



Inflation: Regimes with different persistence

y_t is inflation. $y_t | S_t = j \sim N(\mu_j, \sigma_j^2)$.

- Regime 1: low mean and vol
- Regime 2: high mean and vol, less persistent
- Shaded area if $p_{t|\mathcal{T}}(s_t = 2) > 0.5$



More complicated regime switching models

A natural extension is the K-regime, AR(p) model:

$$y_t = \phi_{0,s_t} + \phi_{1,s_t} y_{t-1} + \dots + \phi_{p,s_t} y_{t-p} + \sigma_{s_t} \varepsilon_t,$$

where $\varepsilon_t \stackrel{i.i.d.}{\sim} N(0, 1)$ and there are K regimes, $S_t \in \{1, \dots, K\}$.

It is possible to have time-varying regime probabilities as function of observable variables

- Using a multinomial logit and exogenous vector X_t to get regime probabilities.

Want more detail?

- Time series analysis by Hamilton
- Or, how about you ask your favorite LLM for further in-depth discussion. It has read all the textbooks....

State Space + Switching Regime Model

SR-SSMs

Canonical Linear-Gaussian SSM

Observation and state:

$$\begin{aligned}y_t &= Zx_t + d + \varepsilon_t, & \varepsilon_t &\sim N(0, H), \\x_t &= Tx_{t-1} + c + R\eta_t, & \eta_t &\sim N(0, Q),\end{aligned}$$

with independence over time and across shocks. Initial $x_0 \sim N(a_{0|0}, P_{0|0})$.

Kalman filter:

$$\begin{aligned}a_{t|t-1} &= Ta_{t-1|t-1} + c, & P_{t|t-1} &= TP_{t-1|t-1}T' + RQR', \\v_t &= y_t - Za_{t|t-1} - d, & F_t &= ZP_{t|t-1}Z' + H, \\K_t &= P_{t|t-1}Z'F_t^{-1}, & a_{t|t} &= a_{t|t-1} + K_tv_t, \\P_{t|t} &= P_{t|t-1} - K_tF_tK_t'.\end{aligned}$$

Log-likelihood: $\ell(\Theta) = \sum_{t=1}^T \log N(v_t; 0, F_t)$.

SR-SSM: Model (K States)

Regime $S_t \in \{1, \dots, K\}$, transition matrix $\Pi = [\pi_{ij}]$.

Conditional on $S_t = j$:

$$\begin{aligned} y_t &= Z_j x_t + d_j + \varepsilon_t, & \varepsilon_t &\sim N(0, H_j), \\ x_t &= T_j x_{t-1} + c_j + R_j \eta_t, & \eta_t &\sim N(0, Q_j). \end{aligned}$$

Initials: $S_0 \sim \pi_0$, $x_0 | S_0 = j \sim N(a_{0|0}^{(j)}, P_{0|0}^{(j)})$.

Parameters: $\Theta = \{(Z_j, d_j, H_j, T_j, c_j, R_j, Q_j)_{j=1}^K, \Pi\}$.

Hamilton–Kim Mixture Kalman Filter

Prediction from $i \rightarrow j$:

$$\begin{aligned} a_{t|t-1}^{(j|i)} &= T_j a_{t-1|t-1}^{(i)} + c_j, \\ P_{t|t-1}^{(j|i)} &= T_j P_{t-1|t-1}^{(i)} T_j' + R_j Q_j R_j'. \end{aligned}$$

Mixing:

$$\begin{aligned} p_{t|t-1}(j) &= \sum_i \pi_{ij} p_{t-1}(i), & \omega_{ij,t} &= \frac{\pi_{ij} p_{t-1}(i)}{p_{t|t-1}(j)}, \\ a_{t|t-1}^{(j)} &= \sum_i \omega_{ij,t} a_{t|t-1}^{(j|i)}, \\ P_{t|t-1}^{(j)} &= \sum_i \omega_{ij,t} \left(P_{t|t-1}^{(j|i)} + (a_{t|t-1}^{(j|i)} - a_{t|t-1}^{(j)}) (a_{t|t-1}^{(j|i)} - a_{t|t-1}^{(j)})' \right). \end{aligned}$$

Update in j :

$$\begin{aligned} v_t^{(j)} &= y_t - Z_j a_{t|t-1}^{(j)} - d_j, & F_t^{(j)} &= Z_j P_{t|t-1}^{(j)} Z_j' + H_j, \\ K_t^{(j)} &= P_{t|t-1}^{(j)} Z_j' (F_t^{(j)})^{-1}, \\ a_{t|t}^{(j)} &= a_{t|t-1}^{(j)} + K_t^{(j)} v_t^{(j)}, & P_{t|t}^{(j)} &= P_{t|t-1}^{(j)} - K_t^{(j)} F_t^{(j)} K_t^{(j)'} . \end{aligned}$$

Regime and State Smoothing

Regime smoothing:

$$p_{t|T}(j) = \frac{p_t(j) \sum_{\ell=1}^K \frac{\pi_{j\ell} \times p_{t+1|\ell}(\ell)}{p_{t+1}(\ell)}}{\sum_{m=1}^K p_t(m) \sum_{\ell=1}^K \frac{\pi_{m\ell} \times p_{t+1|\ell}(\ell)}{p_{t+1}(\ell)}}, \quad t = T-1, \dots, 1.$$

State smoothing (Kim collapse + RTS):

$$\begin{aligned}\omega_{j\ell,t+1}^* &= \frac{\pi_{j\ell} \times p_{t+1|\ell}(\ell) / p_{t+1|t}(\ell)}{\sum_{m=1}^K \pi_{jm} \times p_{t+1|m}(\ell) / p_{t+1|t}(m)}, \quad a_{t+1|T}^{(j,*)} = \sum_{\ell} \omega_{j\ell,t}^* a_{t+1|T}^{(\ell)} \\ P_{t+1|T}^{(j,*)} &= \sum_{\ell} \omega_{j\ell,t}^* \left[P_{t+1|T}^{(\ell)} + \left(a_{t+1|T}^{(\ell)} - a_{t+1|T}^{(j,*)} \right) \left(a_{t+1|T}^{(\ell)} - a_{t+1|T}^{(j,*)} \right)' \right] \\ J_t^{(j)} &= P_{t|t}^{(j)} T_j' (P_{t+1|t}^{(j)})^{-1}, \\ \hat{a}_{t|T}^{(j)} &= a_{t|t}^{(j)} + J_t^{(j)} (\hat{a}_{t+1|T}^{(j,*)} - a_{t+1|t}^{(j)}), \\ \hat{P}_{t|T}^{(j)} &= P_{t|t}^{(j)} + J_t^{(j)} (\hat{P}_{t+1|T}^{(j,*)} - P_{t+1|t}^{(j)}) J_t^{(j)'}.\end{aligned}$$

Collapse:

$$\hat{a}_{t|T} = \sum_j p_{t|T}(j) \hat{a}_{t|T}^{(j)}, \quad \hat{P}_{t|T} = \sum_j p_{t|T}(j) (\hat{P}_{t|T}^{(j)} + (\hat{a}_{t|T}^{(j)} - \hat{a}_{t|T}) (\hat{a}_{t|T}^{(j)} - \hat{a}_{t|T})').$$

2-State Inflation Model

Scalar x_t , inflation y_t , regime $S_t \in \{1, 2\}$:

$$\begin{aligned}y_t &= d_j + x_t + \varepsilon_t, \quad \varepsilon_t \sim N(0, h_j), \\x_t &= c_j + \phi_j x_{t-1} + \eta_t, \quad \eta_t \sim N(0, q_j),\end{aligned}$$

with $j = S_t$. $\Pi = \begin{bmatrix} \pi_{11} & \pi_{12} \\ \pi_{21} & \pi_{22} \end{bmatrix}$.

Filter scalars:

$$a_{t|t-1}^{(j|i)} = c_j + \phi_j a_{t-1|t-1}^{(i)}, \quad P_{t|t-1}^{(j|i)} = \phi_j^2 P_{t-1|t-1}^{(i)} + q_j,$$

$$v_t^{(j)} = y_t - d_j - a_{t|t-1}^{(j)}, \quad F_t^{(j)} = P_{t|t-1}^{(j)} + h_j, \quad K_t^{(j)} = \frac{P_{t|t-1}^{(j)}}{F_t^{(j)}}.$$

Identification and Initialization

Only $\mu_j = d_j + \frac{c_j}{1-\phi_j}$ is identified without restrictions.

- Fix $d_1 = 0$ or impose $d_1 = d_2$.
- Constrain $|\phi_j| < 1$ via $\phi_j = \tanh(\varphi_j)$.
- Often set $h_1 = h_2 = h$.

Initialization options:

- Diffuse: large $P_{0|0}^{(j)}$, $a_{0|0}^{(j)} = 0$, p_0 stationary in Π .
- Regime-stationary: $a_{0|0}^{(j)} = \frac{c_j}{1-\phi_j}$, $P_{0|0}^{(j)} = \frac{q_j}{1-\phi_j^2}$.

Exact Likelihood and Direct MLE

Predictive mixture: $p(y_t | y_{1:t-1}) = \sum_{j=1}^2 N(v_t^{(j)}; 0, F_t^{(j)}) p_{t|t-1}(j).$

Log-likelihood: $\mathcal{L}(\Theta) = \sum_{t=1}^T \log \left(\sum_{j=1}^2 f_j(y_t) p_{t|t-1}(j) \right).$

Use L-BFGS-B or Newton with reparametrizations:

$\phi_j = \tanh \varphi_j, \ h_j = e^{\lambda_j}, \ q_j = e^{\kappa_j}, \ \pi_{12} = \sigma(\alpha), \ \pi_{21} = \sigma(\beta).$

Cost penalty for number of regimes $O(TK^2) = O(4T).$

EM Algorithm

E-step: compute $p_{t|T}(j) = \Pr(S_t = j | y_{1:T})$, $\xi_t(i,j) = \Pr(S_{t-1} = i, S_t = j | Y)$,
and smoothed moments $\bar{x}_t^{(j)}$, $\bar{x}_t^2^{(j)}$, $\bar{x}_t x_{t-1}^{(j)}$.

M-step:

$$\hat{\pi}_{ij} = \frac{\sum_t \xi_t(i,j)}{\sum_t p_{t-1|T}(i)}, \quad \hat{d}_j = \frac{\sum_t p_{t|T}(j) (y_t - \bar{x}_t^{(j)})}{\sum_t p_{t|T}(j)},$$

$$\hat{h}_j = \frac{\sum_t p_{t|T}(j) [(y_t - \hat{d}_j)^2 - 2(y_t - \hat{d}_j)\bar{x}_t^{(j)} + \bar{x}_t^2^{(j)}]}{\sum_t p_{t|T}(j)},$$

$$\begin{bmatrix} \hat{c}_j \\ \hat{\phi}_j \end{bmatrix} = \left(\sum_t p_{t|T}(j) \begin{bmatrix} 1 \\ \bar{x}_{t-1}^{(j)} \end{bmatrix} \begin{bmatrix} 1 & \bar{x}_{t-1}^{(j)} \end{bmatrix} \right)^{-1} \left(\sum_t p_{t|T}(j) \begin{bmatrix} \bar{x}_t^{(j)} \\ \bar{x}_t x_{t-1}^{(j)} \end{bmatrix} \right),$$

$$\hat{q}_j = \frac{1}{\sum_t p_{t|T}(j)} \sum_t p_{t|T}(j) E[(x_t - \hat{c}_j - \hat{\phi}_j x_{t-1})^2 | y_{1:T}, S_t = j].$$

Inflation: 2-state SR-SSM estimation results

	d_1	d_2	c_1	c_2	q_1	q_2
MLE	0.54	0.41	0.06	0.02	0.39	0.07
(s.e.)	(297)	(297)	(6.99)	(0.193)	(0.044)	(0.005)
	ϕ_1	ϕ_2	h	π_{11}	π_{22}	
MLE	0.976	0.999	1.25	0.957	0.984	
(s.e.)	(0.012)	(0.001)	(0.001)	(0.021)	(0.007)	

- Regime 2: low mean, low volatility
- Regime 1: high mean, high volatility
- Both regimes very persistent, but regime 1 less so
 - ▶ Regime 1: Duration 1.9 years; Regime 2: Duration 5.4 years.

Does anything look off to you in the estimates?

Important: you cannot blindly trust LLM output

Despite having read papers on the SR-SSM, the estimation has issues

- ① The observation equation means d_1 and d_2 have enormous standard errors
- ② So does c_1 and to some extent c_2
- ③ Persistence ϕ_2 statistically unit root
- ④ h just one number, not separate volatility for each state in observation equation

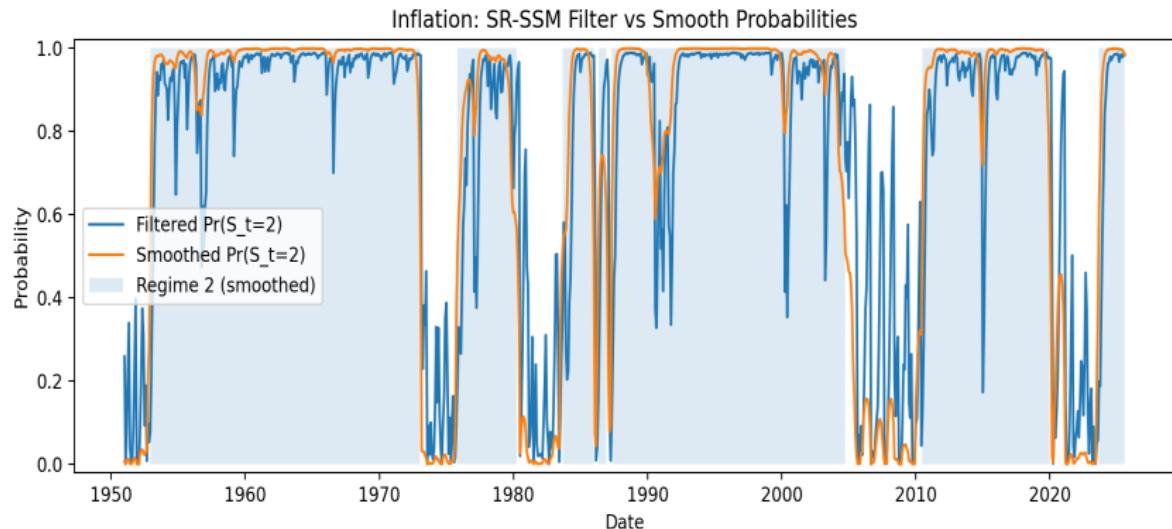
What's going on? Model is overidentified and has close to nonstationary elements.

- Need to choose two parameters of c_1 , c_2 , d_1 , d_2
 - ▶ E.g., constrain $c_1 = c_2 = 0$.
- Perhaps constrain $|\phi_j| < 1 - 2 \times \text{s.e.}(\phi_j)$ for stationarity?
- Perhaps help with starting values for h_1, h_2 set equal to variance of observable, or similar?
- Ask LLM for suggestion for help with better identification of variables in model?

Modeling and estimation is an iterative procedure

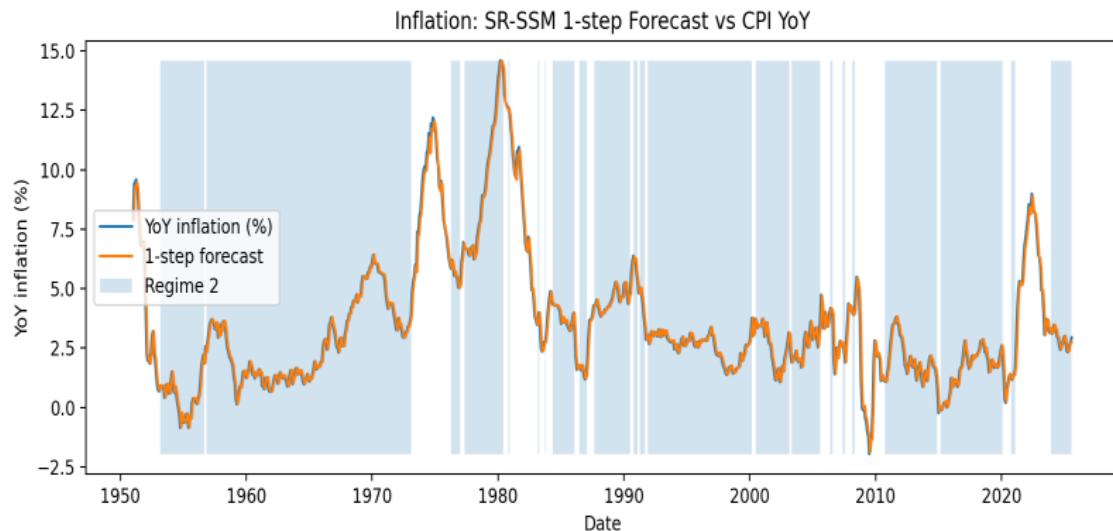
Inflation: Filtering and smoothing probabilities

- Regime 1: low mean and vol
- Regime 2: high mean and vol, less persistent
- Shaded area if $p_{t|T} (s_t = 2) > 0.5$



Inflation: Regimes with different persistence

- Regime 2: low mean and vol
- Regime 1: high mean and vol, less persistent
- Notice much better forecast of inflation than simple SR earlier



Another Nonlinear Version of the SSM

The Unscented Kalman Filter

UKF: Nonlinear SSM

UKF: Unscented Kalman Filter

$$\begin{aligned}y_t &= h(x_t) + \varepsilon_t, \quad \varepsilon_t \sim N(0, R), \\x_t &= f(x_{t-1}) + \eta_t, \quad \eta_t \sim N(0, Q).\end{aligned}$$

Here, $h(\cdot)$ and $f(\cdot)$ are known (you choose) nonlinear functions.

UKF propagates moments via “sigma points” instead of linearization.

UKF: Sigma Points

Given $x \sim N(m, P)$, construct $\{\chi^{(i)}\}_{i=0}^{2n}$ with weights $\{W^{(i)}\}$:

$$\begin{aligned}\chi^{(0)} &= m, \quad \chi^{(i)} = m + [\sqrt{(n+\lambda)P}]_i, \quad i = 1, \dots, n, \\ \chi^{(i)} &= m - [\sqrt{(n+\lambda)P}]_{i-n}, \quad i = n+1, \dots, 2n,\end{aligned}$$

$$\begin{aligned}W_m^{(0)} &= \frac{\lambda}{n+\lambda}, \quad W_c^{(0)} = \frac{\lambda}{n+\lambda} + 1 - \alpha^2 + \beta, \quad W_m^{(i)} = W_c^{(i)} = \frac{1}{2(n+\lambda)}, \\ \lambda &= \alpha^2(n+\kappa) - n.\end{aligned}$$

Main idea: use Kalman filter for each sigma point, then average across sigma points using weights W

UKF: Updates

Time update:

$$\begin{aligned}\chi_{t|t-1}^{(i)} &= f(\chi_{t-1|t-1}^{(i)}), \quad m_{t|t-1} = \sum_i W_m^{(i)} \chi_{t|t-1}^{(i)}, \\ P_{t|t-1} &= \sum_i W_c^{(i)} (\chi_{t|t-1}^{(i)} - m_{t|t-1}) (\chi_{t|t-1}^{(i)} - m_{t|t-1})' + Q.\end{aligned}$$

Measurement update:

$$\begin{aligned}\zeta_t^{(i)} &= h(\chi_{t|t-1}^{(i)}), \quad \hat{y}_t = \sum_i W_m^{(i)} \zeta_t^{(i)}, \\ S_t &= \sum_i W_c^{(i)} (\zeta_t^{(i)} - \hat{y}_t) (\zeta_t^{(i)} - \hat{y}_t)' + R, \\ C_t &= \sum_i W_c^{(i)} (\chi_{t|t-1}^{(i)} - m_{t|t-1}) (\zeta_t^{(i)} - \hat{y}_t)', \\ K_t &= C_t S_t^{-1}, \quad m_{t|t} = m_{t|t-1} + K_t (y_t - \hat{y}_t), \\ P_{t|t} &= P_{t|t-1} - K_t S_t K_t'.\end{aligned}$$

Conclusion

Combining the SSM with non-linear features provides very flexible non-linear modeling tools

Maximum likelihood estimation is standard and has a large well-developed literature that you can draw on

- For instance, you can add a penalty for number of parameters and ridge/lasso type penalties like we did for the logistic regression in Data Analytics in Spring
- In this case, you would need some type of cross-validation to find the optimal constraint for best out-of-sample performance

These SSM-based models are typically relative intuitive and help us understand the underlying mechanisms driving the economy/the process at hand

The downside is that they (largely) assume you know the functional form (e.g., linear + regimes).

- Machine learning methods like neural nets and deep learning are more flexible and aimed at allowing for very flexible nonlinear patterns