

# Lecture 3

## Feed-forward Neural Networks

Lars A. Lochstoer  
UCLA Anderson School of Management

# Overview

## The Universal Approximation Theorem

- Neural nets as universal function approximators

## Deep learning and efficiency: The role of multiple layers

- What does multiple layers buy us?
- Functional compositions

## Feed-forward networks

- Review of structure
- What functions to approximate: SSM, SR, SSM+GARCH
- Approximating state variables using finite-windows
- Estimation: stochastic gradient descent, backpropagation, vanishing gradient problem


## Why Neural Nets?

### The Universal Approximation Theorem

# One-Layer Networks: A Review

Recall the structure of a one-hidden-layer network

- $n_0$  inputs (data)
- $n_1$  units in hidden layer /
- Activation function  $\sigma(\cdot)$  (e.g., sigmoid, ReLU, Tanh)



# One-Layer Networks: Mathematical Representation

Data space:  $x \in \mathbb{R}^{n_0}$ . Target:  $f : K \rightarrow \mathbb{R}$  continuous on compact  $K \subset \mathbb{R}^{n_0}$ .

Single-hidden-layer network with  $n_1$  units:

$$N_{n_1}(x) = \sum_{j=1}^{n_1} \beta_j \sigma(w_j^\top x + b_j) + \beta_0,$$

parameters  $\theta = (\{\beta_j\}_{j=0}^{n_1}, \{w_j, b_j\}_{j=1}^{n_1})$ . Activation function  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ .

Approximation criterion:  $\|f - N_{n_1}\|_\infty = \sup_{x \in K} |f(x) - N_{n_1}(x)|$ .

Goal: for any  $\varepsilon > 0$ , find  $m$  and  $\theta$  with  $\|f - N_{n_1}\|_\infty < \varepsilon$ .

# Universal Approximation Theorem

## Theorem

Let  $K \subset \mathbb{R}^{n_0}$  compact and  $C(K)$  continuous real-valued functions on  $K$ .

If  $\sigma$  is continuous and nonpolynomial (e.g., sigmoid, tanh, ReLU), then the set

$$\mathcal{N} = \left\{ \sum_{j=1}^{n_1} \beta_j \sigma(w_j^\top x + b_j) + \beta_0 : n_1 \in \mathbb{N} \right\}$$

is dense in  $C(K)$  under  $\|\cdot\|_\infty$ .

- “Dense in  $C(K)$ ”: neural net can approximate any continuous function on  $K$  to arbitrary accuracy!

Equivalently,  $\forall f \in C(K)$ ,  $\forall \varepsilon > 0$ ,  $\exists N_{n_1} \in \mathcal{N}$  with  $\|f - N_{n_1}\|_\infty < \varepsilon$ .

Remarks: density does not imply uniqueness or parameter identifiability; rates of convergence depend on smoothness of  $f$  and choice of  $\sigma$ .

## Proof sketch: The ReLU case

ReLU can create a continuous piecewise-linear approximation to a continuous function.

- Easiest to see using a specific example

Let  $f(x) = \sin(x)$ ,  $x \in [0, 2\pi]$ . Choose uniform knots  
 $0 = x_0 < x_1 < \dots < x_{n_1} = 2\pi$ , spacing  $h = 2\pi/n_1$ .



Linear spline  $S_h$  defined by node values  $y_j = f(x_j)$ : for  $x \in [x_j, x_{j+1}]$ ,

$$S_h(x) = \frac{x_{j+1} - x}{h} y_j + \frac{x - x_j}{h} y_{j+1}.$$

- Ie, pick points on true function, linearly interpolate between these points.
- Next slide shows this and hints that with more points (units) we get less approximation error!

# Linear interpolation of continuous function

Intuitively (and it's mathematical fact), the approximation error becomes arbitrarily small as the number of ReLU units goes to infinity



# Why Deep Learning?

## Functional Approximations and Efficiency

# What does multiple hidden layers buy us?

Efficiency.

Functional compositions:  $(f \circ g)(x) = f(g(x))$

- Layers act like functional compositions
- Thus, if the function you are trying to approximate can be expressed as a functional composition, multiple hidden layers is likely more efficient than a single hidden layer


Efficiency: smaller out-of-sample forecast error (e.g., expected mean squared error)

- Recall two types of error: variance and bias
- Multiple layers often have fewer parameters when calibrated to the same approximation error (ie. same bias) as a one-layer network
  - ▶ Thus, less estimation error and more efficient

## Simple example of multilayer efficiency

Consider a tent map (function):

$$T(x) = \max(0, 1 - 2|x - 1/2|) \quad \text{on } x \in [0, 1]$$




Can be fit with 3 ReLu units

- Think long 1 call with  $K = 0$ , short two calls with  $K = 0.5$ , long one call with  $K = 1$

## Simple example: functional composition

Next, consider a case where function is a composition of the same function:



$$T^{(2)} = (T \circ T)(x) = T(T(x))$$



## Simple example: functional composition (cont'd)

General  $L$ -time composition (again, on same function which is special case \_)

$$T^{(L)} = \underbrace{(T \circ T \circ \cdots \circ T)}_{L \text{ times}}(x)$$



- Original function has 2 linear pieces,  $L$  layer composition has  $2^L$  linear pieces
  - ▶ Exponential growth in complexity it allows for!

## Simple example: fitting with ReLU

One-layer ReLU:

$$f(x) = c_0 + c_1 x + \sum_{i=1}^{n_1} a_i \max(0, x - b_i)$$

- Thus,  $2n_1 + 2$  parameters, where  $n_1$  is the number of units in the layer
- You basically need one unit per kink in the function
- You need intercept and slope ( $c_0$  and  $c_1$ ) to fit average level and first slope from left (possible to add another unit instead of the  $c_1x$  term, but this is less efficient as it introduces 2 parameters instead of 1)

Multi-layer approach

- Repeating triangles: need 2 ReLU units per layer

## Simple example: fitting with ReLU (cont'd)

With  $L = 10$  there are  $2^{10} = 1024$  linear pieces and 1023 kinks (with endpoints not counted)

- A one-hidden-layer network thus needs 1023 ReLU units;  $2n_1 + 2 = 2048$  parameters

Multilayer approach: 10 layers, 2 ReLU units per layer, plus final output layer

- Layer 1: two units – 2 weights  $w_i$  and intercepts  $b_i$ ;
- Layer 2-10: two units with two inputs from prior layer: 4 weights  $w_i$  and 2 intercepts  $b_i$ ;
- Final output, linear average: 2 weights and 1 intercept

In sum: 61 parameters with 10-layer network, 2048 with 1-layer

- Clearly  $61 < 2048$ , which is the source of efficiency gain!

## Common functions that use multiple layers

Certain common functions are more efficiently approximated by multiple hidden layers.



- $\exp(\cdot)$  and  $\ln(\cdot)$  are prominent examples

Curvature means many units required in single layer

- As previous example shows, can get in the order of  $n^L$  (often written  $O(n^L)$ ) breakpoints with  $n$  units per layer, where  $L$  is number of layers
  - ▶ The ReLU naturally gives rise to a linear spline approximation and many break points are needed for close approximations
- Plots on next slide show efficiency gain as a function of approximation error  $\varepsilon$

# Common functions: deep learning efficiency

- Think of the size of the model as related to the number of parameters needed to achieve a certain approximation error
- As we go further to the right on the x-axis, the approximation error becomes smaller
- The one-layer network (sometimes refer to as 1-D) explodes in size, while the multi-layer approach does not.



## Form of activation function

The above discussion assumed ReLU activation function  $\sigma(\cdot)$

- Convenient as piecewise linear is easy to understand

Other forms are also common

- Which is best depends on function we are trying to approximate
- Typically, we use cross-validation and MSE to choose number of layers and type of activation function
- Still, there are choices we can make beyond this approach if we know something about the function we want to approximate

Other common activation functions include:

$$\text{sigmoid: } \frac{e^u}{1 + e^u},$$

$$\tanh: \frac{e^u - e^{-u}}{e^u + e^{-u}}.$$

What types of functions are we trying to approximate?

# Benchmarks

Before we start fitting a bunch of data and look at out-of-sample MSEs....

- Let's think about some properties of the type of functions we are looking to approximate

Three examples:

- ① SSM: simple linear Gaussian case
- ② SR: simple 2-state switching model
- ③ SSM-GARCH: SSM model with time-varying variance of shocks

What do optimal forecasts look like in these models, and how can FFNNs (feed-forward neural nets) approximate this?

## Linear state-space model forecasts

Standard state-space models have forecasts of the form:

$$\begin{aligned}E_t(y_{t+1}) &= Za_{t+1|t} + d \\&= \phi_0 + \sum_{j=1}^{\infty} \phi_j y_{t+1-j}.\end{aligned}$$

(here prior is not present as we have implicitly assumed stationary dynamics and an infinite history)

- That is, the forecast as a  $AR(\infty)$  (or  $VAR(\infty)$  in vector case)

ReLU can fit a linear function with two units:

$$a + bx = a + b \max(0, x) - b \max(0, -x)$$

- Note: input vector to FFNN should be *all historical data* given the infinite lag requirement
- In practice, choose window length  $J$ , which implicitly sets  $\phi_j = 0$  for  $j > J$ 
  - Decent approximation if system not too persistent
  - Thus: at time  $t$  feed the FFNN  $y_t, y_{t-1}, \dots, y_{t-J+1}$

## Switching regime model forecasts

Consider simple 2-state model  $y_t|s_t \sim N(\mu_{s_t}, \sigma_{s_t}^2)$  with transition probability matrix  $\Pi$

- From previous lecture note:

$$E_t(y_{t+1}) = \mathbf{p}'_t \Pi \mu,$$

where  $\mu' = [\mu_1 \quad \mu_2]$  and  $\mathbf{p}'_t = [p_t(s_t = 1) \quad p_t(s_t = 2)]$

So we need to express beliefs as a function of the history of observables

- Recall, update for belief that  $s_t = j$ :

$$\begin{aligned} p_t(j) &= \frac{1}{c_t} f_j(y_t) p_{t|t-1}(j) \\ &= \frac{1}{c_t} f_j(y_t) \sum_{i=1}^2 \pi_{ij} p_{t-1}(i), \end{aligned}$$

where  $c_t = \sum_{m=1}^2 f_m(y_t) p_{t|t-1}(m)$  is the normalizing constant so probabilities sum to 1 and  $\pi_{ij} = \Pr(S_t = j | S_{t-1} = i)$

## Forecast function based on observables

We can iterate backwards on the updating equation:

$$p_t(j) = \frac{1}{c_t} f_j(y_t) \sum_{i=1}^2 \pi_{ij} \frac{1}{c_{t-1}} f_i(y_{t-1}) p_{t-1|t-2}(i)$$

- Keep going and we get a function that is a sum of exponentials of lagged  $y^t$ : multiple layers will be efficient!

Note: if  $f_j(y_t) > f_i(y_t)$  the data is saying state  $j$  was more likely

- Thus, lagged sequence of data tells us if we should weight  $\mu_1$  or  $\mu_2$  more in our current forecast
- Again, need full history of  $y^t$  as input vector, but in practice choose a window of length  $J$  where  $J$  depends on the persistence of the system

## SSM-GARCH model forecasts

**Model:**  $y_t = Zx_t + d + \varepsilon_t$ ,  $x_t = \phi x_{t-1} + \eta_t$ , with  $\varepsilon_t \sim N(0, h_t)$  and  $\eta_t \sim N(0, q)$

- Note that observation equation has time-varying volatility of noise,  $h_t$

Optimal one-step forecast:  $E_t(y_{t+1}) = Z\phi a_{t|t} + d$

- The forecast admits an  $AR(\infty)$ -type representation

$$E_t(y_{t+1}) = c_t + \sum_{j=0}^{\infty} \psi_{j,t} y_{t-j},$$

where  $c_t$  and  $\psi_{j,t}$  are functions of the history of the variances of the noise terms,  $h^t$ , as well as the model parameters ( $Z, \phi, q$ )

- Interpretation: larger past  $h_{t-j}$  reduces the effective weight  $\psi_{j,t}$  assigned to  $y_{t-j}$ . The vector  $h^t$  therefore acts like an “attention” vector that down-weights noisier observations

# Understanding FFNNs

## A Monte-Carlo Perspective

## Notation for deep feed-forward network

Input  $x$  is  $n_0 \times 1$  vector of data.

Output of each hidden layer  $l = \{1, \dots, L\}$

$$\mathbf{a}^{[l]} = \sigma^{[l]} \left( \mathbf{W}^{[l]} \mathbf{a}^{[l-1]} + \mathbf{b}^{[l]} \right),$$

where the activation function is applied elementwise.

- Layer  $l$  has  $n_l$  units. Thus,  $\mathbf{a}^{[l-1]}$  is  $n_{l-1} \times 1$ ,  $\mathbf{W}^{[l]}$  is  $n_l \times n_{l-1}$ ,  $\mathbf{a}^{[l]}$  and  $\mathbf{b}^{[l]}$  are  $n_l \times 1$
- It's also useful to define:  $\mathbf{h}^{[l]} = \mathbf{W}^{[l]} \mathbf{a}^{[l-1]} + \mathbf{b}^{[l]}$

Output layer:

$$\mathbf{y} = \mathbf{U} \mathbf{a}^{[L]},$$

where  $\mathbf{U}$  is  $n_{L+1} \times n_L$ . Here  $n_{L+1} \geq 1$ .


- Often, if  $n_{L+1} > 1$ , we are looking for a probability distribution where we take the additional step

$$\tilde{y}^{(k)} = \text{softmax } \left( y^{(k)} \right) = \frac{\exp \left( y^{(k)} \right)}{\sum_{j=1}^{n_{L+1}} \exp \left( y^{(j)} \right)},$$

where superscript  $(j)$  denotes row  $j$

# Canonical 2-Layer feed-forward network

(a) with bias node; (b) bias node included in  $x$  vector



In this plot, the output layer has many variables (e.g., for a probability distribution)

## Estimating the weights

Often Stochastic Gradient Descent is used to estimate the function

- Use random subset of data (a mini-batch of length  $b$ ) to update gradient/parameter guess
- An “Epoch” is one full pass over training set (number of iterations per Epoch  $T/b$  if  $T$  is size of sample)
- Typical Epoch settings is 20-300 depending on case
- Early Stopping means stop when objective function fails to improve

Update rule:

$$\theta_{t+1} = \theta_t - \eta g_{B_t}(\theta_t),$$

where  $g_{B_t}(\theta_t)$  is the mini-batch gradient for iteration  $t$  and  $\eta$  is the learning rate.

- $\theta_t$  is a vector with all the parameters

## Overview of Monte Carlo exercise

- **Goal:** learn a feed-forward NN that maps a short window of past observations to the one-step-ahead forecast.
- **Data:** 2-state Gaussian switching-regime (HMM “emissions”) with asymmetric transitions.
- **Training:** only uses  $y$  (and optionally a noisy proxy  $z$ ); loss is MSE on realized  $y_{t+1}$ .
- **Evaluation:** compare NN forecast against *true* conditional mean from HMM filter and an AR(1) baseline.

# Monte Carlo: true model (HMM emissions)

**State process:**  $s_t \in \{0, 1\}$ , Markov with

$$\Pi = \begin{bmatrix} 0.95 & 0.05 \\ 0.20 & 0.80 \end{bmatrix}, \quad \mathbf{p}_{t|t-1} = \Pi \mathbf{p}_{t-1}.$$

**Observation:**

$$y_t \mid s_t = j \sim \mathcal{N}(\mu_j, \sigma^2), \quad \mu = (-1, 1), \quad \sigma = 1.$$

**Filtering and forecast:** with posterior  $\mathbf{p}_t = \Pr(s_t = \cdot \mid y_{1:t})$ ,

$$m_t \equiv \mathbb{E}[y_{t+1} \mid y_{1:t}] = \mathbf{p}'_t \Pi \mu, \quad p_t(i) \propto p_{t|t-1}(i) f_i(y_t).$$

**Irreducible risk vs realized  $y_{t+1}$ :**

$$\text{Var}(y_{t+1} \mid y_{1:t}) = \sigma^2 + (\mu_1 - \mu_0)^2 \mathbf{p}'_{t+1|t} (1 - \mathbf{p}_{t+1|t}).$$

# Neural net specification

**Inputs** (window  $J = 2$ ):  $X_t = [y_{t-1}, y_t]$ .


**Architecture:**

$$X \xrightarrow{\text{Aff+ReLU}} a_1 \xrightarrow{\text{Aff+ReLU}} a_2 \xrightarrow{\text{Aff}} \hat{y}_{t+1|t}.$$

Hidden sizes  $n_1 = n_2 = 64$ ; ReLU activations; linear output.

**Loss (training on realized  $y_{t+1}$ ):**

$$\mathcal{L} = \frac{1}{T} \sum_t (\hat{y}_t - y_t^{\text{obs}})^2 + \frac{\lambda}{2} \sum \|W\|_F^2.$$



# Estimation of Feed-forward Neural Net

**Input:** at each  $t$ , input  $y_t, y_{t-1}, \dots, y_{t-J+1}$

**Numerical optimizer:** Typically, Stochastic Gradient Descent and MSE loss function plus L1 (Lasso) and/or L2 (Ridge) constraint

## Stochastic Gradient Descent:

- Gradient based method (derivative of objective function wrt parameters in network)
  - ▶ Use backpropagation to compute gradient (or derivative vector, or score)
- Stochastic: Choose a random mini-batch (random sub-set of data) for each iteration of gradient descent
- Epoch: is one complete pass through the data.
  - ▶ Multiple epochs typically chosen to fit model until some early stopping criterion (e.g., parameters stops updating meaningfully) is achieved

## Backpropagation: goal

**Goal.** Efficiently compute the gradient  $\nabla_{\theta}\mathcal{L}$  of a scalar loss  $\mathcal{L}$  w.r.t. millions of parameters  $\theta$  so we can *learn* via gradient methods.

### Why it is needed.

- Directly differentiating the full composition is costly and error-prone.
- Finite differences need  $O(\#\theta)$  forward passes and are noisy.
- Backpropagation uses the *chain rule* organized as dynamic programming to get all partials in time proportional to one forward/backward sweep.

Recall:

$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f(g(x))}{\partial g(x)} \frac{\partial g(x)}{\partial x}$$

## Backpropagation: core idea

**Core idea.** Treat the network as a computation graph of local functions. For a layer output  $a^{(\ell)}$  feeding later nodes, the sensitivities of the loss with respect to node values are given by (using Chain Rule directly here)

$$\underbrace{\frac{\partial \mathcal{L}}{\partial h^{(\ell)}}}_{\text{upstream "blame"}} = \sum_{k \in \text{children of } \ell} \underbrace{\frac{\partial \mathcal{L}}{\partial h^{(k)}}}_{\text{already known local Jacobian}} \underbrace{\frac{\partial h^{(k)}}{\partial h^{(\ell)}}}_{\cdot}$$

Propagate these sensitivities *backward* from the loss to inputs; then obtain parameter gradients by one more local multiplication:

$$\nabla_{W^{(\ell)}} \mathcal{L} = (a^{(\ell-1)})^\top \underbrace{\frac{\partial \mathcal{L}}{\partial h^{(\ell)}}}_{\Delta^{(\ell)}}, \quad \nabla_{b^{(\ell)}} \mathcal{L} = \mathbf{1}^\top \Delta^{(\ell)}.$$

- To understand last equation, elementwise we have

$$\frac{\partial \mathcal{L}}{\partial W_{ij}^{(\ell)}} = \sum_r \frac{\partial \mathcal{L}}{\partial h_r^{(\ell)}} \frac{\partial h_r^{(\ell)}}{\partial W_{ij}} = \frac{\partial \mathcal{L}}{\partial h^{(\ell)}} a_j^{(\ell-1)} \quad \text{where } \frac{\partial h_r^{(\ell)}}{\partial W_{ij}} = \mathbf{1}_{\{r=i\}} a_j^{(\ell-1)}$$

# Backpropagation: purpose

## What it accomplishes.

- *Efficiency:* Computes  $\nabla_{\theta}\mathcal{L}$  in  $O(\text{forward cost})$ , independent of  $\#\theta$  up to constants.
- *Correct attribution:* Each weight learns how a small change would reduce  $\mathcal{L}$  given the current batch.
- *Scalability:* Works for any differentiable graph (ReLU, normalization, residuals), enabling deep nets.

**Econometric analogy.** Backprop is a *reverse recursion of scores*: local “scores” (sensitivities) are propagated backward, like filtering vs. smoothing forward computes predictions, backward attributes error to earlier transformations.

# Backpropagation: algorithm, notation, and tips

**Define:**  $\|W_1\|_F^2 = \text{trace}(W_1^\top W_1) = \sum_i \sum_j (W_1)_{ij}^2$ .

## Mini-batch Stochastic Gradient Descent (one step).

① *Forward:* compute  $h_1, a_1, h_2, a_2, \hat{y}$ .

② *Loss:*  $\mathcal{L} = \frac{1}{N} \sum (\hat{y} - y)^2 + \frac{\lambda}{2} \sum \|W\|_F^2$ .

③ *Backward:*

$$\Delta = \frac{2}{N}(\hat{y} - y), \quad \Delta_{a_2} = \Delta W_y^\top, \quad \Delta_{h_2} = \Delta_{a_2} \odot \mathbb{I}[h_2 > 0],$$

$$\Delta_{a_1} = \Delta_{h_2} W_2^\top, \quad \Delta_{h_1} = \Delta_{a_1} \odot \mathbb{I}[h_1 > 0].$$

Accumulate  $\nabla_W, \nabla_b$  as on the derivation slide.

④ *Update:*  $W \leftarrow W - \eta \nabla_W, \quad b \leftarrow b - \eta \nabla_b$ .

## Notes.

- ReLU mask  $\mathbb{I}[h > 0]$  zeros gradients where units are inactive.
- Weight decay: add  $\lambda W$  to  $\nabla_W$  only (not to biases).
- Standardize  $X$ ; pick small  $\eta$ ; clip gradients if needed (to bound step size)
- All formulas are vectorized over the batch; avoid loops.

# Backpropagation: derivation for the MLP

**Network (mini-batch size  $N$ ).**

$$\begin{aligned} X &\in \mathbb{R}^{N \times J} \quad (\text{window length } J), \\ h_1 &= XW_1 + \mathbf{1}b_1, \quad a_1 = \sigma(h_1), \\ h_2 &= a_1 W_2 + \mathbf{1}b_2, \quad a_2 = \sigma(h_2), \\ \hat{y} &= a_2 W_y + \mathbf{1}b_y \in \mathbb{R}^{N \times 1}, \\ \sigma(u) &= \max\{0, u\} \quad (\text{ReLU}). \end{aligned}$$

**Loss (MSE +  $L_2$ ).**

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2 + \frac{\lambda}{2} \left( \|W_1\|_F^2 + \|W_2\|_F^2 + \|W_y\|_F^2 \right).$$

# Backpropagation: derivation for the MLP (cont'd)

**Output gradient.**

$$\frac{\partial \mathcal{L}}{\partial \hat{y}} = \frac{2}{N}(\hat{y} - y) \equiv \Delta.$$

**Readout layer.**

$$\nabla_{W_y} = a_2^\top \Delta + \lambda W_y, \quad \nabla_{b_y} = \mathbf{1}^\top \Delta, \quad \Delta_{a_2} = \Delta W_y^\top.$$

**Second hidden layer.**

$$\Delta_{h_2} = \Delta_{a_2} \odot \mathbb{I}[h_2 > 0], \quad \nabla_{W_2} = a_1^\top \Delta_{h_2} + \lambda W_2, \quad \nabla_{b_2} = \mathbf{1}^\top \Delta_{h_2}, \quad \Delta_{a_1} = \Delta_{h_2} W_2^\top$$

**First hidden layer.**

$$\Delta_{h_1} = \Delta_{a_1} \odot \mathbb{I}[h_1 > 0], \quad \nabla_{W_1} = X^\top \Delta_{h_1} + \lambda W_1, \quad \nabla_{b_1} = \mathbf{1}^\top \Delta_{h_1}.$$

**Shapes:**  $W_1 : J \times n_1$ ,  $W_2 : n_1 \times n_2$ ,  $W_y : n_2 \times 1$ .

# Backpropagation and SGD

**Gradients:**

$$\frac{\partial \mathcal{L}}{\partial \hat{y}} = \frac{2}{N}(\hat{y} - y), \quad \frac{\partial \mathcal{L}}{\partial h_2} = \left( \frac{\partial \mathcal{L}}{\partial \hat{y}} W_y^\top \right) \odot \mathbf{1}_{h_2 > 0},$$

$$\nabla_{W_2} = a_1^\top \frac{\partial \mathcal{L}}{\partial h_2} + \lambda W_2, \quad \nabla_{W_1} = X^\top \left( \left( \frac{\partial \mathcal{L}}{\partial h_2} W_2^\top \right) \odot \mathbf{1}_{h_1 > 0} \right) + \lambda W_1.$$

**SGD update:**  $\theta \leftarrow \theta - \eta \nabla_\theta$  on mini-batches.

**Code excerpt** (note code has a and h reversed, sorry!):

```
dy = (2.0/N)*(y_hat - y_true)
dWy = h2.T @ dy + weight_decay*Wy
dh2 = dy @ Wy.T
da2 = dh2 * relu_grad(a2)
dW2 = h1.T @ da2 + weight_decay*W2
...
W2 -= lr*dW2; b2 -= lr*db2; Wy -= lr*dWy; by -= lr*dby
```

# Evaluation protocol

**True forecast (teacher) for comparison only:**

$$m_t = \mathbb{E}[y_{t+1} \mid y_{1:t}] = \mu^\top (\pi_t P), \quad \text{taken from hmm\_filter.}$$

**Alignment:** windows ending at  $t$  map to  $m_t$  and to targets  $y_{t+1}$  at index  $t + 1$ .



**Metrics vs  $m_t$ :**

$$\text{MSE} = \frac{1}{N} \sum (\hat{y} - m)^2, \quad R^2 = 1 - \frac{\sum (\hat{y} - m)^2}{\sum (m - \bar{m})^2}.$$

**Plots:**

- ①  $m_t$  vs NN forecast.
- ②  $m_t$  vs AR(1) forecast.

# Results from SR Monte Carlo exercise



Reading: NN forecast tracks regime-driven shifts in the true mean better than AR(1). Discrepancies remain near switches due to residual mixture uncertainty.

# Results from SR Monte Carlo exercise

## Validation metrics vs true conditional mean $m_t$

Model	MSE	$R^2$
MLP	0.079	0.684
AR(1)	0.237	0.050

*Reading:* NN forecast tracks regime-driven shifts in the true mean better than AR(1) due to capacity for handling nonlinear models.


## Adding a noisy regime proxy

**Observe:**  $z_t = p_t(1) + \eta_t$ ,  $\eta_t \sim \mathcal{N}(0, \sigma_z^2)$ , where  $p_t(1) = \Pr(s_t = 1 \mid y_{1:t})$

- Noisy but informative signal of what state we are in

**Augmented input:**  $[y_{t-1}, y_t, z_{t-1}, z_t]$ .

**Training unchanged:** MSE on  $y_{t+1}$ .



## Results using additional signal

With signal model performs much better

**Validation metrics vs true conditional mean  $m_t$**

Model	MSE	$R^2$
MLP	0.018	0.912
AR(1)	0.250	-0.227

*Reading:* NN forecast tracks regime-driven shifts in the true mean better than AR(1) due to capacity for handling nonlinear models.

# Warning: Vanishing gradients problem

## Setup.

- Depth- $L$  MLP with ReLU/sigmoid/tanh hidden layers:

$$h^{(\ell)} = a^{(\ell-1)} W_\ell + \mathbf{1} b_\ell, \quad a^{(\ell)} = \sigma(h^{(\ell)}), \quad \ell = 1, \dots, L.$$

## Backprop chain.

- For loss  $\mathcal{L}$  and layer  $\ell$ :

$$\frac{\partial \mathcal{L}}{\partial W_\ell} = (a^{(\ell-1)})^\top \Delta^{(\ell)}, \quad \Delta^{(\ell)} = \Delta^{(\ell+1)} W_{\ell+1}^\top \odot \sigma'(h^{(\ell)}).$$

Unrolling from the output ( $\Delta^{(L+1)}$  known):

$$\Delta^{(\ell)} = \Delta^{(L+1)} \left( \prod_{k=\ell+1}^L W_k^\top D^{(k-1)} \right), \quad D^{(k)} = \text{diag}(\sigma'(k)).$$

# Vanishing gradients (cont'd)

## Why vanishing.

- If  $\|\cdot\|$  is any sub-multiplicative matrix norm,

$$\|\Delta^{(\ell)}\| \leq \|\Delta^{(L+1)}\| \prod_{k=\ell+1}^L \|W_k\| \|D^{(k-1)}\|.$$

- For sigmoid/tanh,  $|\sigma'(u)| \leq 1/4$  (sigmoid) and  $|\sigma'(u)| \leq 1$  with saturation  $\approx 0$  (tanh). Hence  $\|D^{(k)}\| \ll 1$  often, so the product decays  $\rightarrow 0$  as depth grows.

## Saturation view.

- In saturated regions,  $\sigma'^{(k)} \approx 0 \Rightarrow \Delta^{(\ell)} \approx 0$ . Parameters in early layers receive near-zero updates.

**Plainly:** when you multiply together derivatives that are less than  $|1|$  in a backpropagation chain, they very quickly go to zero, which means the optimizer doesn't find an effect as it blocks learning about the parameter

# Mitigations: ReLU and batch normalization

## ReLU non-vanishing region.

$$\sigma(u) = \max\{0, u\}, \quad \sigma'(u) = \begin{cases} 1, & u > 0 \\ 0, & u < 0 \text{ ("off")} \end{cases}$$

In the active region  $u > 0$ ,  $D^{(k)} = \mathbf{I}$ . Then along active paths,

$$\Delta^{(\ell)} \approx \Delta^{(L+1)} \prod_{k=\ell+1}^L W_k^\top,$$

so gradients do not systematically shrink due to activations.

- *Caveat:* dead ReLUs ( $u < 0$  persistently) yield zero gradients.

## Batch normalization (BN).

- Technique that uses normalization of pre-activations  $h^{(\ell)}$  in a mini-batch to ensure gradient does not vanish

## In sum...

Neural nets are “simply” flexible approximating functions

Feed-forward networks have nice properties

- Estimating network, forecasting, backpropagation can be done in parallel
- I.e., perfect for GPU computing

Downside: FFNN have no endogenous state-variables

- SR-SSM have sufficient statistics (state and latent variable beliefs)
  - ▶ Yield low dimensional representation
  - ▶ But, comes from a sequential structure so not parallelizable

HW 3 you will work with FFNNs

- Next lecture: Recurrent Neural Nets (RNNs) and other network structures that allow endogenous state-variables working in a sequential fashion